

Week 05 Lectures

Projection on Primary Key

1/73

No duplicates, so the above approaches are not required.

Method:

```

bR = nPages(Rel)
for i in 0 .. bR-1 {
    P = read page i
    for j in 0 .. nTuples(P) {
        T = getTuple(P, j)
        T' = mkTuple(pk, T)
        if (outBuf is full) write and clear
        append T' to outBuf
    }
}
if (nTuples(outBuf) > 0) write
  
```

Get all pages

Index-only Projection

2/73

Can do projection without accessing data file iff ...

- relation is indexed on (A_1, A_2, \dots, A_n) (indexes described later)
- projected attributes are a prefix of (A_1, A_2, \dots, A_n)

two common of indexing:

1. ID PageID
2. Id_1 ~ Id_n : page_01
Id_n ~ Id_2n : page_02

Basic idea:

- scan through index file (which is already sorted on attributes)
- duplicates are already adjacent in index, so easy to skip

Cost analysis ...

- index has b_i pages (where $b_i \ll b_R$)
- Cost = b_i reads + b_{Out} writes

Comparison of Projection Methods

3/73

Difficult to compare, since they make different assumptions:

- index-only: needs an appropriate index
- hash-based: needs buffers and good hash functions
- sort-based: needs only buffers \Rightarrow use as default

Best case scenario for each (assuming $n+1$ in-memory buffers):

- index-only: $b_i + b_{Out} \ll b_R + b_{Out}$
- hash-based: $b_R + 2.b_P + b_{Out}$
- sort-based: $b_R + b_T + 2.b_T \cdot \text{ceil}(\log_n b_0) + b_T + b_{Out}$

We normally omit b_{Out} since each method produces the same result

Projection in PostgreSQL

4/73

Code for projection forms part of execution iterators:

- `backend/executor/execQual.c`

Functions involved with projection:

- `ExecProject(projInfo, ...)` ... extracts projected data

- `check_sql_fn_retval(...)` ... makes new tuple via TargetList
- `ExecStoreTuple(newTuple,...)` ... save tuple in buffer

plus many many others ...

Implementing Selection

Varieties of Selection

6/73

Selection: `select * from R where C`

- filters a subset of tuples from one relation *R*
- based on a condition *C* on the attribute values

We consider three distinct styles of selection:

- 1-d (**one dimensional**) (condition uses only 1 attribute)
- *n*-d (**multi-dimensional**) (condition uses >1 attribute)
- similarity (approximate matching, with ranking)

Each style has several possible file-structures/techniques.

... Varieties of Selection

7/73

Examples of different selection types:

- *one*: `select * from R where id = 1234`
- *pmr*: `select * from R where age=65 (1-d)`

`select * from R where age=65 and gender='m' (n-d)`
- *rng*: `select * from R where age≥18 and age≤21 (1-d)`

`select * from R where age between 18 and 21 (n-d)`
`and height between 160 and 190`
note: rng = range
- *sim*: `select * from R where name like '%oo%'`

`select * from Images where similar to SampleImage`

Exercise 1: Query Types

8/73

Using the relation:

All primary key has an index

```
create table Courses (
  id          integer primary key,
  code        char(8),  -- e.g. 'COMP9315'
  title       text,     -- e.g. 'Computing 1'
  year        integer,  -- e.g. 2000..2016
  convenor    integer references Staff(id),
  constraint once_per_year unique (code,year)
);
```

give examples of each of the following query types:

1. a 1-d one query, an *n*-d one query
2. a 1-d *pmr* query, an *n*-d *pmr* query
3. a 1-d range query, an *n*-d range query

Suggest how many solutions each might produce ...

Implementing Select Efficiently

9/73

Two basic approaches:

- physical arrangement of tuples
 - sorting (search strategy)
 - hashing (static, dynamic, n-dimensional)
- additional indexing information
 - index files (primary, secondary, trees)
 - signatures (superimposed, disjoint)

Our analysis assumes 1 input buffer available for each relation.

If more buffers are available, most methods benefit.

Heap Files

Note: this is **not** "heap" as in the top-to-bottom ordered tree.
It means simply an unordered collection of tuples in a file.

Selection in Heaps

11/73

For all selection queries, the only possible strategy is:


```
// select * from R where C
for each page P in file of relation R {
    for each tuple t in page P {
        if (t satisfies C)
            add tuple t to result set
    }
}
```

i.e. linear scan through file searching for matching tuples

... Selection in Heaps

12/73

The heap is scanned from the first to the last page:



$$Cost_{range} = Cost_{pmr} = b$$

If we know that only one tuple matches the query (one query),
a simple optimisation is to stop the scan once that tuple is found.

$$Cost_{one}: \quad Best = 1 \quad Average = b/2 \quad Worst = b$$

Insertion in Heaps

13/73

Insertion: new tuple is appended to file (in last page).

```
rel = openRelation("R", READ|WRITE);
pid = nPages(rel)-1;
get_page(rel, pid, buf);
if (size(newTup) > size(buf))
    { deal with oversize tuple }
else {
    if (!hasSpace(buf, newTup))
        { pid++; nPages(rel)++; clear(buf); }
    insert_record(buf, newTup);
    put_page(rel, pid, buf);
}
```

Always insert tuple at last pages

$$Cost_{insert} = 1_r + 1_w$$

... Insertion in Heaps

14/73

Alternative strategy:

- find any page from R with enough space
- preferably a page already loaded into memory buffer

PostgreSQL's strategy:


- use last updated page of R in buffer pool
- otherwise, search buffer pool for page with enough space
- assisted by free space map (FSM) associated with each table
- for details: [backend/access/heap/{heapam.c,hio.c}](#)

... Insertion in Heaps

15/73

Dealing with oversize tuple t :

```
for i in 1 .. nAttr(t) {
    if (t[i] not oversized) continue
    off = appendToFile(ovsf, t[i])
    t[i] = (OVERSIZE, off)
}
insert into buf as before
```



... Insertion in Heaps

16/73

PostgreSQL's tuple insertion:

```
heap_insert(Relation relation, // relation desc
            HeapTuple newtup,   // new tuple data
            CommandId cid, ...) // SQL statement
```

- finds page which has enough free space for newtup
- ensures page loaded into buffer pool and locked
- copies tuple data into page buffer, sets xmin, etc.
- marks buffer as dirty
- writes details of insertion into transaction log
- returns OID of new tuple if relation has OIDs

Deletion in Heaps

17/73

SQL: delete from R where Condition

Implementation of deletion:

```
rel = openRelation("R", READ|WRITE);
for (p = 0; p < nPages(rel); p++) {
    get_page(rel, p, buf);
    ndels = 0;
    for (i = 0; i < nTuples(buf); i++) {
        tup = get_record(buf, i);
        if (tup satisfies Condition)
            { ndels++; delete_record(buf, i); }
    }
}
```

```

    if (ndels > 0) put_page(rel, p, buf);
    if (ndels > 0 && unique) break;
}

```

Exercise 2: Cost of Deletion in Heaps

18/73

Consider the following queries ...

```

delete from Employees where id = 12345 -- one
delete from Employees where dept = 'Marketing' -- pmr
delete from Employees where 40 <= age and age < 50 -- range

```

Show how each will be executed and estimate the cost, assuming:

- $b = 100$, $b_{q2} = 3$, $b_{q3} = 20$

State any other assumptions.

... Deletion in Heaps

19/73

PostgreSQL tuple deletion:

```

heap_delete(Relation relation, // relation desc
            ItemPointer tid, ..., // tupleID
            CommandId cid, ...) // SQL statement

```

- gets page containing tuple tid into buffer pool and locks it
- sets flags, commandID and $xmax$ in tuple; dirties buffer
- writes indication of deletion to transaction log

Vacuuming eventually compacts space in each page.

Updates in Heaps == delete + insert

20/73

SQL: update R set $F = val$ where Condition

Analysis for updates is similar to that for deletion

- scan all pages
- replace any updated tuples (within each page)
- write affected pages to disk

$$Cost_{update} = b_r + b_{qw}$$

Complication: new tuple larger than old version (too big for page)

Solution: delete, re-organise free space, then insert

... Updates in Heaps

21/73

PostgreSQL tuple update:

```

heap_update(Relation relation, // relation desc
            ItemPointer otid, // old tupleID
            HeapTuple newtup, ..., // new tuple data
            CommandId cid, ...) // SQL statement

```

- essentially does $delete(otid)$, then $insert(newtup)$
- also, sets old tuple's $ctid$ field to reference new tuple
- can also update-in-place if no referencing transactions

Heaps in PostgreSQL

22/73

PostgreSQL stores all table data in heap files (by default).

Typically there are also associated index files.

If a file is more useful in some other form:

- PostgreSQL may make a transformed copy during query execution
- programmer can set it via `create index...using hash`

Heap file implementation: <src/backend/access/heap>

... Heaps in PostgreSQL

23/73

PostgreSQL "heap file" may use multiple physical files

- files are named after the OID of the corresponding table
- first data file is called simply `OID`
- if size exceeds 1GB, create a fork called `OID.1`
- add more forks as data size grows (one fork for each 1GB)
- other files:
 - free space map (`OID_fsm`), visibility map (`OID_vm`)
 - optionally, TOAST file (if table has varlen attributes)
- for details: Chapter 68 in PostgreSQL v12 documentation

Sorted Files


Sorted Files

25/73

Records stored in file in order of some field k (the sort key).

Makes searching more efficient; makes insertion less efficient


E.g. assume $c = 4$



... Sorted Files

26/73

In order to mitigate insertion costs, use overflow pages.



Total number of overflow pages = b_{ov} .

Average overflow chain length = $O_v = b_{ov} / b$.

Bucket = data page + its overflow page(s)

keep records of the max and min value of each page(bucket, including the overflow page)

Selection in Sorted Files

27/73

For one queries on sort key, use binary search.

```
// select * from R where k = val (sorted on R.k)
lo = 0; hi = b-1
while (lo <= hi) {
    mid = (lo+hi) / 2; // int division with truncation
    (tup, loVal, hiVal) = searchBucket(f, mid, x, val);
    if (tup != NULL) return tup;
    else if (val < loVal) hi = mid - 1;
    else if (val > hiVal) lo = mid + 1;
    else return NOT_FOUND;
}
return NOT_FOUND;
```

where *f* is file for relation, *mid*, *lo*, *hi* are page indexes,
k is a field/attr, *val*, *loVal*, *hiVal* are values for *k*

... Selection in Sorted Files

28/73

Search a page and its overflow chain for a key value

```
searchBucket(f, p, k, val)
{
    buf = getPage(f, p);
    (tup, min, max) = searchPage(buf, k, val, +INF, -INF)
    if (tup != NULL) return (tup, min, max);
    ovf = openOvFile(f);
    ovp = overflow(buf);
    while (tup == NULL && ovp != NO_PAGE) {
        buf = getPage(ovf, ovp);
        (tup, min, max) = searchPage(buf, k, val, min, max)
        ovp = overflow(buf);
    }
    return (tup, min, max);
}
```

Assumes each page contains index of next page in Ov chain

Note: `getPage(f, pid) = { read_page(relOf(f), pid, buf); return buf; }`

... Selection in Sorted Files

29/73

Search within a page for key; also find min/max key values

```
searchPage(buf, k, val, min, max)
{
    res = NULL;
    for (i = 0; i < nTuples(buf); i++) {
        tup = getTuple(buf, i);
        if (tup.k == val) res = tup;
        if (tup.k < min) min = tup.k;
        if (tup.k > max) max = tup.k;
    }
    return (res, min, max);
}
```

... Selection in Sorted Files

30/73

The above method treats each bucket like a single large page.

Cases:

- best: find tuple in first data page we read
- worst: full binary search, and not found
 - examine $\log_2 b$ data pages
 - plus examine all of their overflow pages
- average: examine some data pages + their overflow pages


$Cost_{one} : \text{Best} = 1 \quad \text{Worst} = \log_2 b + b_{ov}$

Average case cost analysis needs assumptions (e.g. data distribution)

Exercise 3: Searching in Sorted File

31/73

Consider this sorted file with overflows ($b=5, c=4$):



Compute the cost for answering each of the following:

- `select * from R where k = 24` ¹ ???
- `select * from R where k = 3` ³
- `select * from R where k = 14`
- `select max(k) from R` ¹

Exercise 4: Optimising Sorted-file Search

32/73

The `searchBucket(f, p, k, val)` function requires:

- read the p^{th} page from data file
- scan it to find a match and min/max k values in page
- while no match, repeat the above for each overflow page
- if we find a match in any page, return it
- otherwise, remember min/max over all pages in bucket

Suggest an optimisation that would improve `searchBucket()` performance for most buckets.


... Selection in Sorted Files

33/73

For pmr query, on non-unique attribute k , where file is sorted on k

- tuples containing k may span several pages

E.g. `select * from R where k = 2`



Begin by locating a page p containing $k=val$ (as for one query).

Scan backwards and forwards from p to find matches.

Thus, $Cost_{pmr} = Cost_{one} + (b_q - 1) \cdot (1 + Ov)$ What is b_q here? Here $b_q = 3$


... Selection in Sorted Files

34/73

For range queries on unique sort key (e.g. primary key):

- use binary search to find lower bound
- read sequentially until reach upper bound

E.g. `select * from R where k >= 5 and k <= 13`



$$Cost_{range} = Cost_{one} + (b_q - 1) \cdot (1 + Ov)$$

... Selection in Sorted Files

35/73

For range queries on non-unique sort key, similar method to pmr.

- binary search to find lower bound
- then go backwards to start of run
- then go forwards to last occurrence of upper-bound

E.g. select * from R where k >= 2 and k <= 6



$$Cost_{range} = Cost_{one} + (b_q - 1) \cdot (1 + Ov)$$

... Selection in Sorted Files

36/73

So far, have assumed query condition involves sort key k.

But what about select * from R where j = 100.0 ?

If condition contains attribute j, not the sort key

- file is unlikely to be sorted by j as well
- sortedness gives no searching benefits

$Cost_{one}$, $Cost_{range}$, $Cost_{pmr}$ as for heap files

Insertion into Sorted Files


37/73

Insertion approach:

- find appropriate page for tuple (via binary search)
- if page not full, insert into page
- otherwise, insert into next overflow page with space

Thus, $Cost_{insert} = Cost_{one} + \delta_w$ (where $\delta_w = 1$ or 2)

Consider insertions of k=33, k=25, k=99 into:



Deletion from Sorted Files

38/73

E.g. delete from R where k = 2

Deletion strategy:

- find matching tuple(s)
- mark them as deleted

Cost depends on *selectivity* of selection condition

Recall: selectivity determines b_q (# pages with matches)


Thus, $Cost_{delete} = Cost_{select} + b_{qw}$

Hashed Files

Hashing

40/73

Basic idea: use key value to compute page address of tuple.



e.g. tuple with key = v is stored in page i

Requires: hash function $h(v)$ that maps $KeyDomain \rightarrow [0..b-1]$.

- hashing converts key value (any type) into integer value
- integer value is then mapped to page index
- note: can view integer value as a bit-string

... Hashing

41/73

PostgreSQL hash function (simplified):

```
Datum hash_any(unsigned char *k, int keylen)
{
    uint32 a, b, c, len, *ka = (uint32 *)k;
    /* Set up the internal state */
    len = keylen;
    a = b = c = 0x9e3779b9+1len+3923095;
    /* handle most of the key */
    while (len >= 12) {
        a += ka[0]; b += ka[1]; c += ka[2];
        mix(a, b, c);
        ka += 3; len -= 12;
    }
    ... collect data from remaining bytes into a,b,c ...
    mix(a, b, c);
    return UInt32GetDatum(c);
}
```

See [backend/access/hash/hashfunc.c](#) for details (incl mix())

... Hashing

42/73

hash_any() gives hash value as 32-bit quantity (uint32).

Two ways to map raw hash value into a page address:

- if $b = 2^k$, bitwise AND with k low-order bits set to one

```
uint32 hashToPageNum(uint32 hval) {
    uint32 mask = 0xFFFFFFFF;
    return (hval & (mask >> (32-k)));
}
```

- otherwise, use *mod* to produce value in range $0..b-1$

```
uint32 hashToPageNum(uint32 hval) {
    return (hval % b);
}
```

Hashing Performance

43/73

Aims:

- distribute tuples evenly amongst buckets
- have most buckets nearly full (attempt to minimise wasted space)

Note: if data distribution not uniform, address distribution can't be uniform.

Best case: every bucket contains same number of tuples.

Worst case: every tuple hashes to same bucket.

Average case: some buckets have more tuples than others.

Use overflow pages to handle "overfull" buckets (cf. sorted files)

All tuples in each bucket must have same hash value.

... Hashing Performance

44/73

r: total number of tuples

Two important measures for hash files:

- load factor: $L = r/bc$
- average **overflow** chain length: $Ov = b_{ov}/b$

Three cases for distribution of tuples in a hashed file:

Case	L	Ov
Best	≈ 1	0
Worst	$\gg 1$	**
Average	< 1	$0 < Ov < 1$

(** performance is same as Heap File)

To achieve average case, aim for $0.75 \leq L \leq 0.9$.

Selection with Hashing

45/73

Select via hashing on unique key *k* (*one*)

```
// select * from R where k = val
P = getPageViaHash(val, R)
for each tuple t in page P {
    if (t.k == val) return t
}
for each overflow page Q of P {
    for each tuple t in page Q {
        if (t.k == val) return t
    }
}
```

$Cost_{one}$: Best = 1, Avg = $1 + Ov/2$ Worst = $1 + max(OvLen)$

... Selection with Hashing

46/73

Working out which page, given a key ...

```

Page getPageViaHash(Value key, Reln R)
{
    Page p; // eventually references a buffer
    uint32 h = hash_any(key, len(key));
    PageID pid = h % nPages(R);
    p = getPage(dataFile(R), pid);
    return p;
}

```

... Selection with Hashing

47/73

Select via hashing on non-unique hash key nk (pmr)

```

// select * from R where nk = val
P = getPageViaHash(val, R)
for each tuple t in page P {
    if (t.nk == val) add t to results
}
for each overflow page Q of P {
    for each tuple t in page Q {
        if (t.nk == val) add t to results
    }
}
return results

```

$$Cost_{pmr} = 1 + Ov$$

If Ov is small (e.g. 0 or 1), very good retrieval cost

... Selection with Hashing

48/73

Hashing does not help with *range* queries** ...

$$Cost_{range} = b + b_{Ov}$$

Selection on attribute j which is not hash key ...

$$Cost_{one}, Cost_{range}, Cost_{pmr} = b + b_{Ov}$$

** unless the hash function is order-preserving (and most aren't)

Insertion with Hashing

49/73

Insertion uses similar process to *one* queries.

```

// insert tuple t with key=val into rel R
P = getPageViaHash(val, R)
if room in page P {
    insert t into P; return
}
for each overflow page Q of P {
    if room in page Q {
        insert t into Q; return
    }
}
add new overflow page Q
link Q to previous page
insert t into Q

```

$$Cost_{insert}: \text{ Best: } 1_r + 1_w \quad \text{ Worst: } 1 + \max(Ov, Len)_r + 2_w$$

Need to add a pointer for the new-added page

Exercise 5: Insertion into Static Hashed File

50/73

Consider a file with $b=4$, $c=3$, $d=2$, $h(x) = \text{bits}(d, \text{hash}(x))$

Insert tuples in alpha order with the following keys and hashes:

<i>k</i>	<i>hash(k)</i>	<i>k</i>	<i>hash(k)</i>	<i>k</i>	<i>hash(k)</i>	<i>k</i>	<i>hash(k)</i>
a	10001	g	00000	m	11001	s	01110
b	11010	h	00000	n	01000	t	10011
c	01111	i	10010	o	00110	u	00010
d	01111	j	10110	p	11101	v	11111
e	01100	k	00101	q	00010	w	10000
f	00010	l	00101	r	00000	x	00111

The hash values are the 5 lower-order bits from the full 32-bit hash.

Deletion with Hashing

51/73

Similar performance to select on non-unique key:

```
// delete from R where k = val
// f = data file ... ovf = overflow file
P = getPageViaHash(val,R)
ndel = delTuples(P,k,val)
if (ndel > 0) putPage(f,P,P.pid)
for each overflow page Q of P {
    ndel = delTuples(Q,k,val)
    if (ndel > 0) putPage(ovf,Q.pid)
}
```

Extra cost over select is cost of writing back modified pages.

Method works for both unique and non-unique hash keys.

Problem with Hashing...

52/73

So far, discussion of hashing has assumed a fixed file size (*b*).

What size file to use?

- the size we need right now (performance degrades as file overflows)
- **the maximum size we might ever need** (significant waste of space)

Problem: change file size \Rightarrow change hash function \Rightarrow rebuild file

Methods for hashing with files whose size changes:

- **extendible hashing, dynamic hashing** (need a directory, no overflows)
- **linear hashing** (expands file "systematically", no directory, has overflows)

Flexible Hashing

53/73

All flexible hashing methods ...

- treat hash as 32-bit bit-string
- adjust hashing by using more/less bits

Start with hash function to convert value to bit-string:

```
uint32 hash(unsigned char *val)
```

Require a function to extract *d* bits from bit-string:

```
uint32 bits(int d, uint32 val)
```

Use result of `bits()` as page address.

Exercise 6: Bit Manipulation

54/73

1. Write a function to display `uint32` values as `01010110...`

```
char *showBits(uint32 val, char *buf);
```

Analogous to `gets()` (assumes supplied buffer is large enough)

2. Write a function to extract the d bits of a `uint32`

```
uint32 bits(int d, uint32 val);
```

If $d > 0$, gives low-order bits; if $d < 0$, gives high-order bits

... Flexible Hashing

55/73

Important concept for flexible hashing: *splitting*

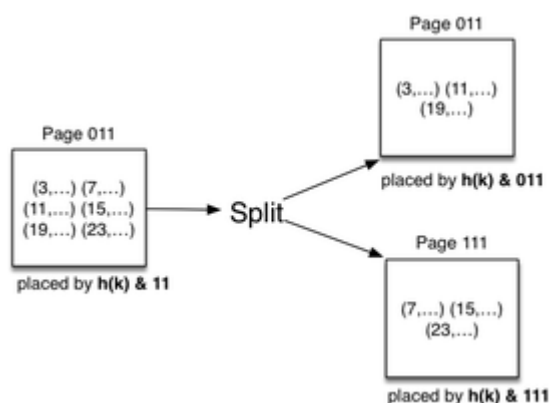
- consider one page (all tuples have same hash value)
- recompute page numbers by considering one extra bit
- if current page is 101, new pages have hashes 0101 and 1101
- some tuples stay in page 0101 (was 101)
- some tuples move to page 1101 (new page)
- also, rehash any tuples in overflow pages of page 101

Result: expandable data file, never requiring a complete file rebuild

... Flexible Hashing

56/73

Example of splitting:



Tuples only show key value; assume $h(val) = val$

Linear Hashing

57/73

File organisation:

- file of primary data pages
- file of overflow data pages
- a register called the *split pointer* (*sp*)

Uses systematic method of growing data file ...

- hash function "adapts" to changing address range
- systematic splitting controls length of overflow chains

Advantage: does *not* require auxiliary storage for a directory


Disadvantage: requires overflow pages (don't split on full pages)

... Linear Hashing

58/73

File grows linearly (one page at a time, at regular intervals).

Has "phases" of expansion; over each phase, b doubles.



Selection with Lin.Hashing

59/73

If $b=2^d$, the file behaves exactly like standard hashing.

Use d bits of hash to compute page address.


```
// select * from R where k = val
h = hash(val);
P = bits(d,h); // lower-order bits
for each tuple t in page P
    and its overflow pages {
        if (t.k == val) return t;
    }
```

Average $Cost_{one} = 1 + Ov$

... Selection with Lin.Hashing

60/73

If $b \neq 2^d$, treat different parts of the file differently.



Parts A and C are treated as if part of a file of size 2^{d+1} .

Part B is treated as if part of a file of size 2^d .

Part D does not yet exist (tuples in B may eventually move into it).

... Selection with Lin.Hashing

61/73

Modified search algorithm:

```
// select * from R where k = val
h = hash(val);
pid = bits(d,h);
if (pid < sp) { pid = bits(d+1,h); }
P = getPage(f, pid)
for each tuple t in page P
    and its overflow pages {
```

Require a function to extract d bits from bit-string:

```
unit32 bits(int d, uint32 val)
```


```

    if (t.k == val) return R;
}

```

File Expansion with Lin.Hashing

62/73



Insertion with Lin.Hashing

63/73

Abstract view:

```

pid = bits(d,hash(val));
if (pid < sp) pid = bits(d+1,hash(val));
// bucket P = page P + its overflow pages
P = getPage(f,pid)
for each page Q in bucket P {
    if (space in Q) {
        insert tuple into Q
        break
    }
}
if (no insertion) {
    add new overflow page to bucket P
    insert tuple into new page
}
if (need to split) {
    partition tuples from bucket sp
    into buckets sp and sp+2^d
    sp++;
    if (sp == 2^d) { d++; sp = 0; }
}

```

Splitting

64/73

How to decide that we "need to split"?

Two approaches to triggering a split:

- split every time a tuple is inserted into full page
- split when load factor reaches threshold (every k inserts)

Note: always split page sp , even if not full or "current"


Systematic splitting like this ...

- eventually reduces length of every overflow chain
- helps to maintain short average overflow chain length

... Splitting

65/73

Splitting process for page $sp=01$:



Exercise 7: Insertion into Linear Hashed File

66/73

4 pages, each page has 3 tuples

Consider a file with $b=4$, $c=3$, $d=2$, $sp=0$, $hash(x)$ as above

Insert tuples in alpha order with the following keys and hashes:

k	$hash(k)$	k	$hash(k)$	k	$hash(k)$	k	$hash(k)$
a	10001	g	00000	m	11001	s	01110
b	11010	h	00000	n	01000	t	10011
c	01111	i	10010	o	00110	u	00010
d	01111	j	10110	p	11101	v	11111
e	01100	k	00101	q	00010	w	10000
f	00010	l	00101	r	00000	x	00111

The hash values are the 5 lower-order bits from the full 32-bit hash.

... Splitting

67/73

Splitting algorithm:

```
// partition tuples between two buckets
newp = sp + 2^d; oldp = sp;
for all tuples t in P[oldp] and its overflows {
    p = bits(d+1, hash(t.k));
    if (p == newp)
        add tuple t to bucket[newp]
    else
        add tuple t to bucket[oldp]
}
sp++;
if (sp == 2^d) { d++; sp = 0; }
```

Insertion Cost

68/73

If no split required, cost same as for standard hashing:

$Cost_{insert}$: Best: $1_r + 1_w$ Avg: $(1+Ov)_r + 1_w$ Worst: $(1+max(Ov))_r + 2_w$

If split occurs, incur $Cost_{insert}$ plus cost of splitting:

- read page sp (plus all of its overflow pages)
- write page sp (and its new overflow pages)
- write page $sp+2^d$ (and its new overflow pages)

On average, $Cost_{split} = (1+Ov)_r + (2+Ov)_w$

Deletion with Lin.Hashing

69/73

Deletion is similar to ordinary static hash file.

But might wish to contract file when enough tuples removed.

Rationale: r shrinks, b stays large \Rightarrow wasted space.

Method:

- remove last bucket in data file (contracts linearly).
- merge tuples from bucket with its buddy page (using $d-1$ hash bits)

Hash Files in PostgreSQL

70/73

PostgreSQL uses linear hashing on tables which have been:

`create index I_x on R using hash (k);`

Hash file implementation: **backend/access/hash**

- **hashfunc.c** ... a family of hash functions
- **hashinsert.c** ... insert, with overflows
- **hashpage.c** ... utilities + splitting
- **hashsearch.c** ... iterator for hash files

Based on "A New Hashing Package for Unix", Margo Seltzer, Winter Usenix 1991

... Hash Files in PostgreSQL

71/73

PostgreSQL uses slightly different file organisation ...

- has a single file containing main and overflow pages
- has groups of main pages of size 2^n
- in between groups, arbitrary number of overflow pages
- maintains collection of group pointers in header page
- each group pointer indicates start of main page group

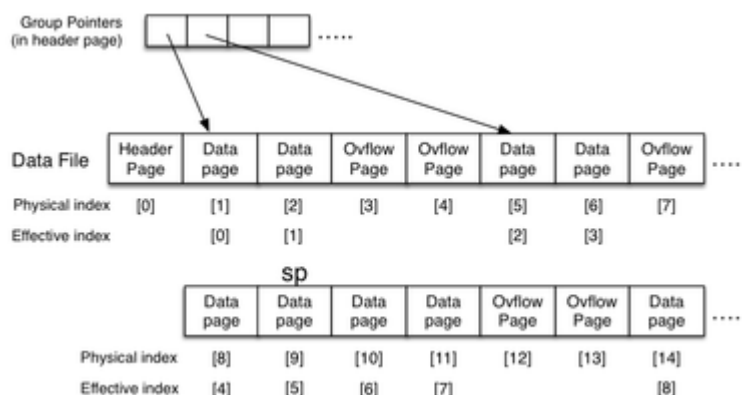
If overflow pages become empty, add to free list and re-use.

Confusingly, PostgreSQL calls "group pointers" as "split pointers"

... Hash Files in PostgreSQL

72/73

PostgreSQL hash file structure:



... Hash Files in PostgreSQL

73/73

Converting bucket # to page address:

```
// which page is primary page of bucket
uint bucket_to_page(headerp, B) {
    uint *splits = headerp->hashm_spares;
    uint chunk, base, offset, lg2(uint);
    chunk = (B<2) ? 0 : lg2(B+1)-1;
    base = splits[chunk];
    offset = (B<2) ? B : B-(1<<chunk);
    return (base + offset);
}
// returns ceil(log_2(n))
int lg2(uint n) {
    int i, v;
    for (i = 0, v = 1; v < n; v <= 1) i++;
    return i;
}
```

Produced: 17 Mar 2020