COMP 3331/9331: Computer Networks and Applications

Week 8

Network Layer: Control Plane (Routing)

Chapter 5: Section 5.1 – 5.2, 5.6

Network layer, control plane: outline

5.1 introduction

5.2 routing protocols

- link state
- distance vector
- hierarchical routing

Self study

5.6 ICMP: The Internet Control Message Protocol

Network-layer functions

Recall: two network-layer functions:

- *forwarding:* move packets from router' s input to appropriate router output

data plane

- *routing:* determine route taken by packets from source to destination

control plane

Two approaches to structuring network control plane:

- per-router control (traditional)
- logically centralized control (software defined networking)

Per-router control plane

Individual routing algorithm components *in each and every router* interact with each other in control plane to compute forwarding tables

distributed control

Distributibutely work out the forwarding table

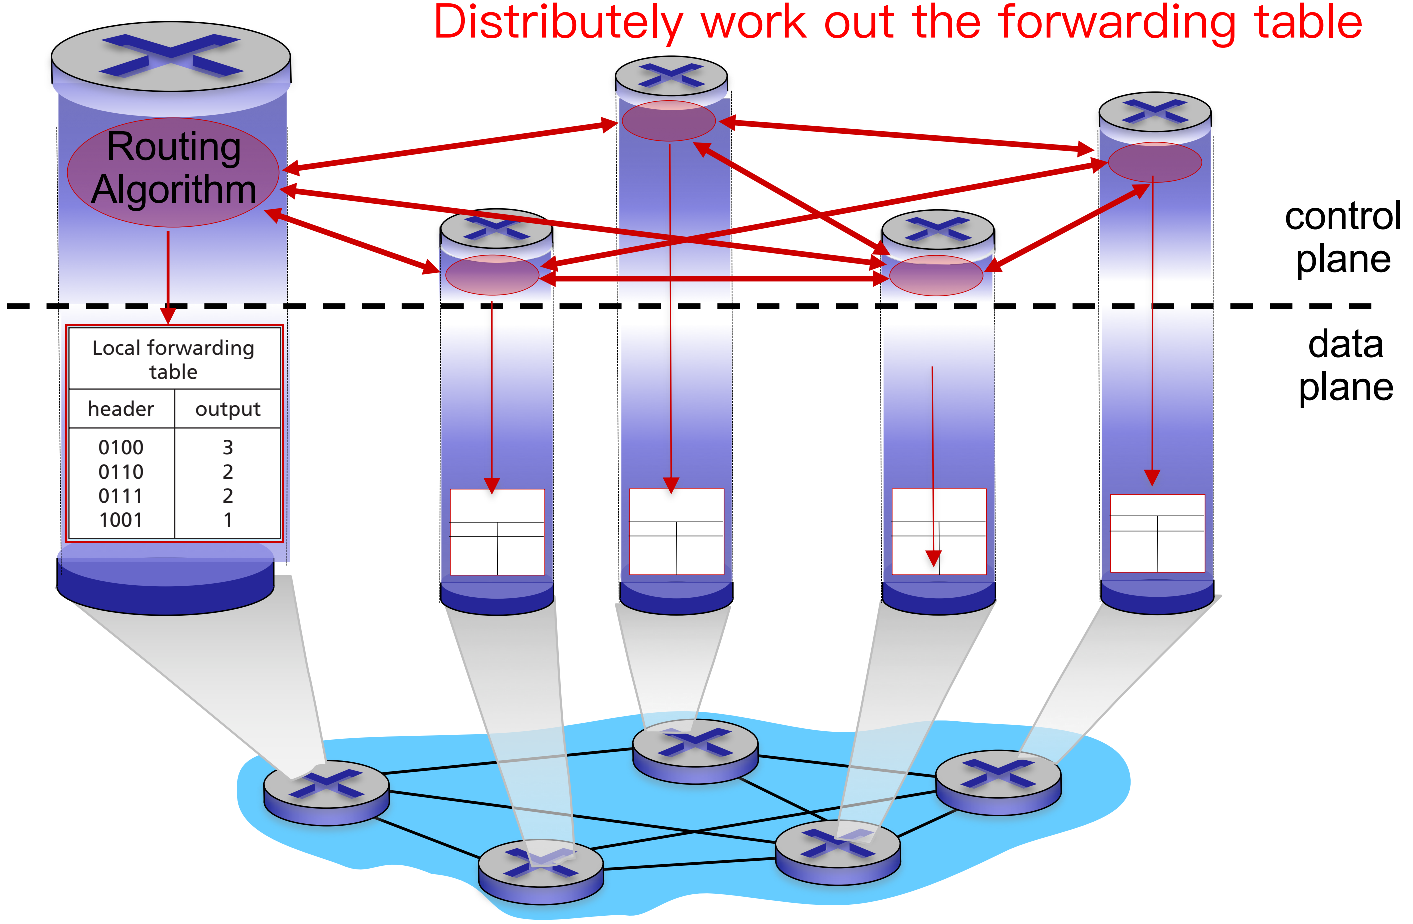

header	output
0100	3
0110	2
0111	2
1001	1

Logically centralized control plane

A distinct (typically remote) controller interacts with local control agents (CAs) in routers to compute forwarding tables

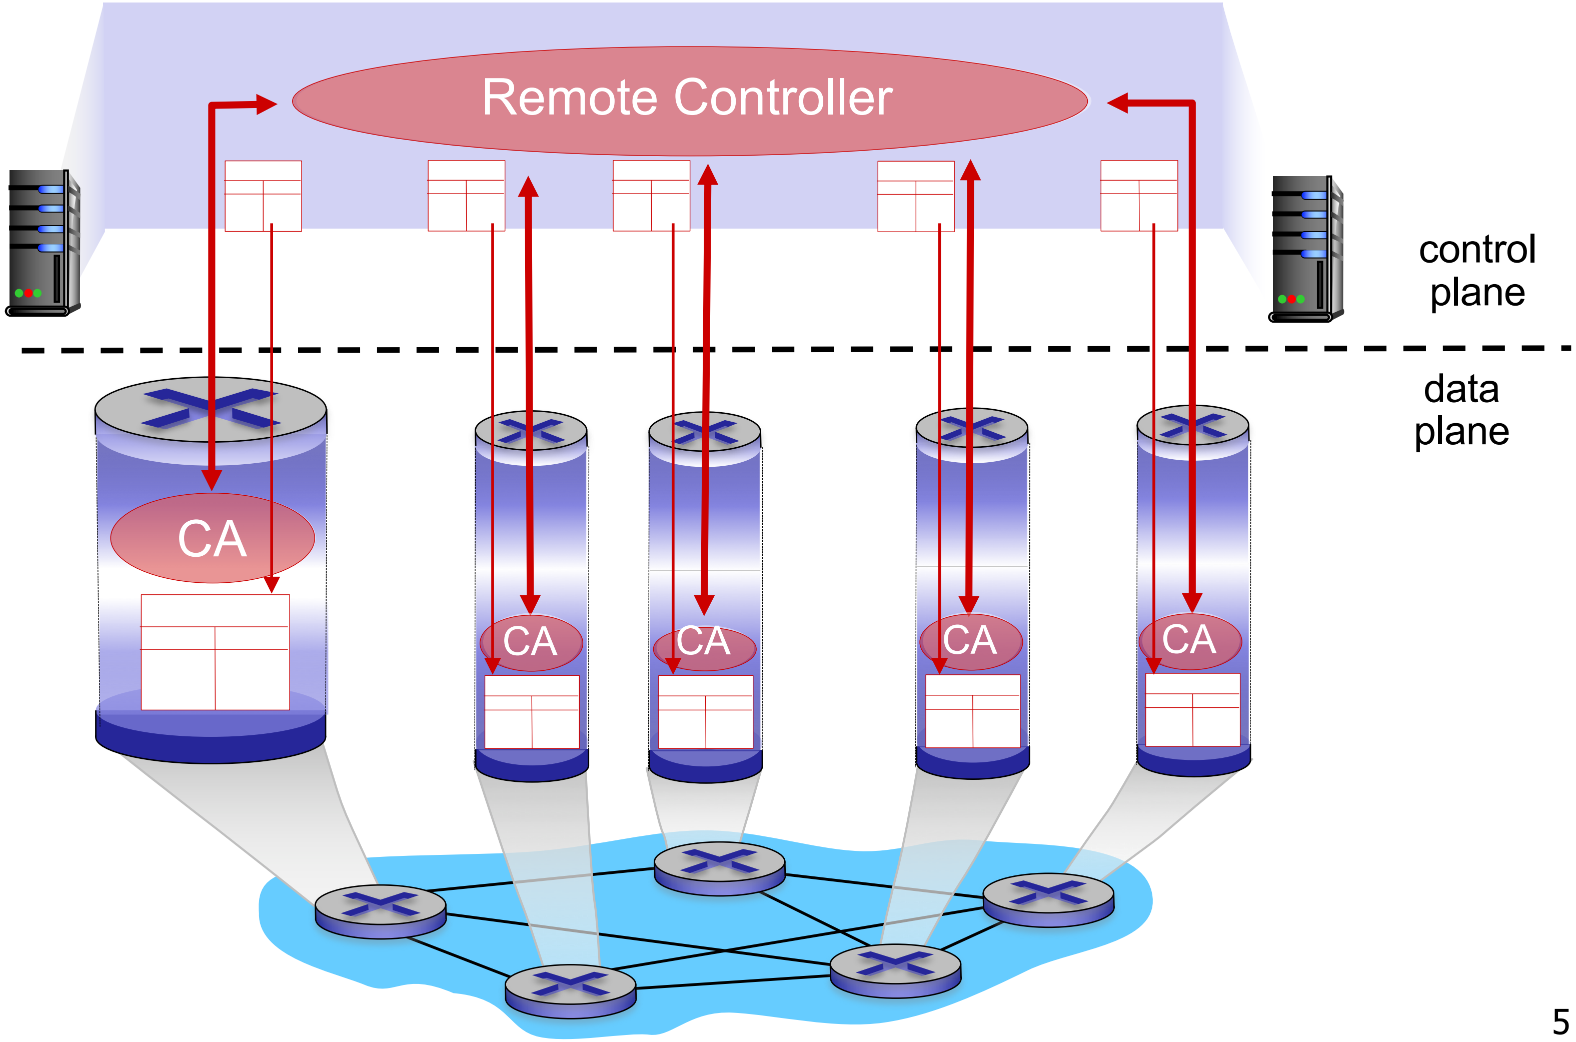

Network layer, control plane: outline

5.1 introduction

5.2 routing protocols

- link state
- distance vector
- Hierarchical routing

5.6 ICMP: The Internet Control Message Protocol

Network under the control of a administration

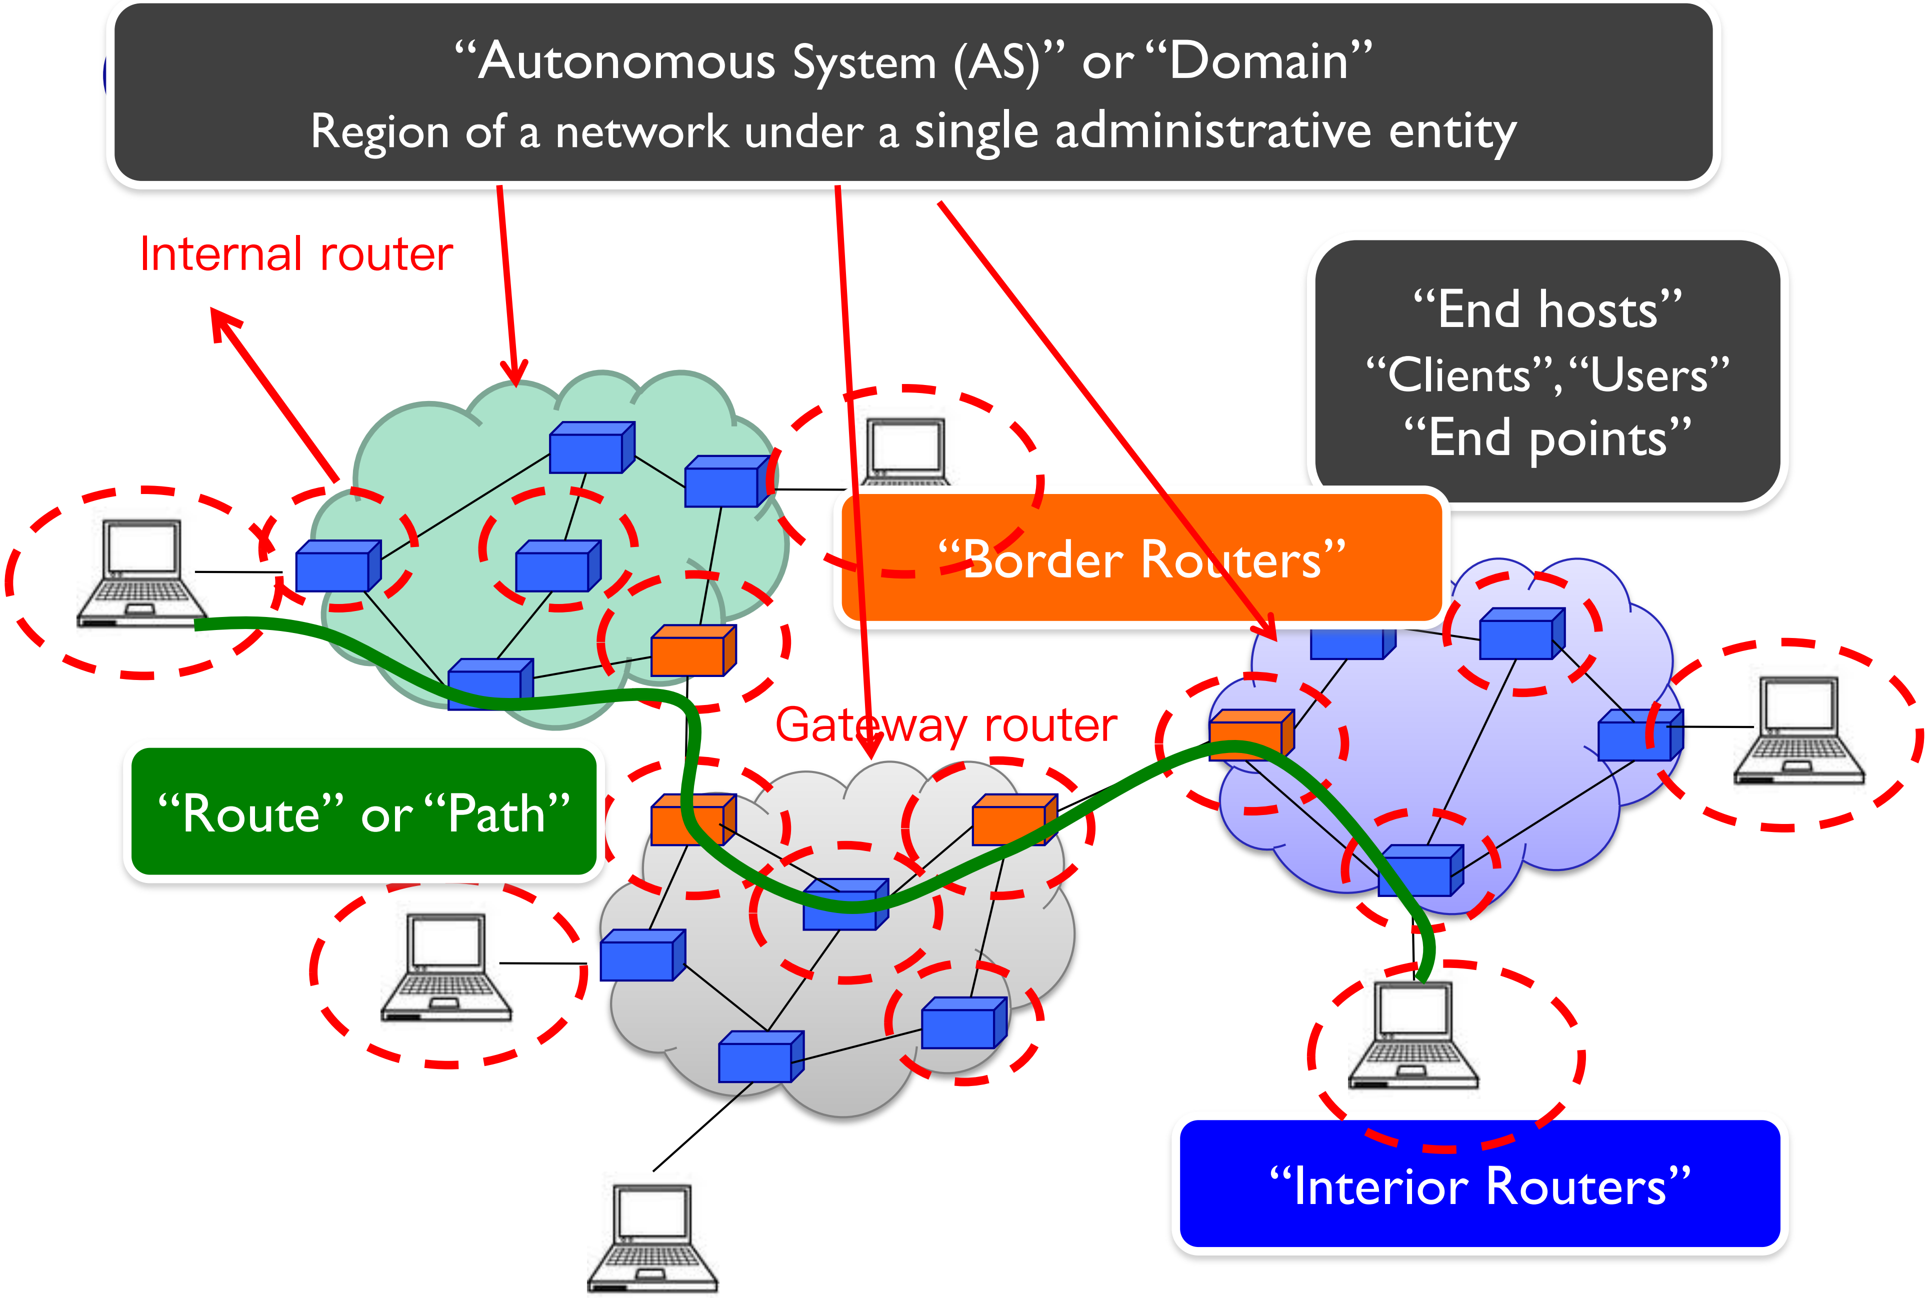
“Autonomous System (AS)” or “Domain”
Region of a network under a single administrative entity

Internal router

“End hosts”
“Clients”, “Users”
“End points”

“Border Routers”

Gateway router

“Route” or “Path”

“Interior Routers”

Internet Routing

- Internet Routing works at two levels

- Each AS runs an intra-domain routing protocol that establishes routes within its domain
  - AS -- region of network under a single administrative entity
  - Link State, e.g., Open Shortest Path First (OSPF)
  - Distance Vector, e.g., Routing Information Protocol (RIP)

- ASes participate in an inter-domain routing protocol that establishes routes between domains
  - Path Vector, e.g., Border Gateway Protocol (BGP)

Graph abstraction

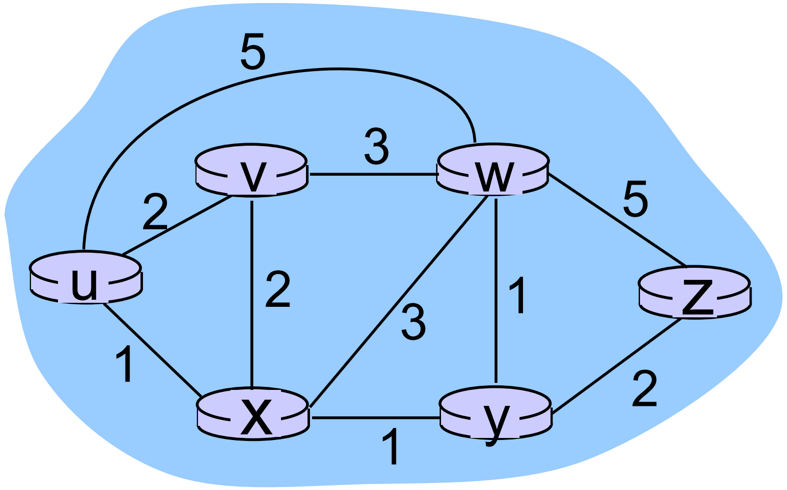

graph: G = (N,E)

N = set of routers = { u, v, w, x, y, z }

E = set of links ={ (u,v), (u,x), (u,w), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) }

It can be ‘point–to–point’ link
Or network

Graph abstraction: costs

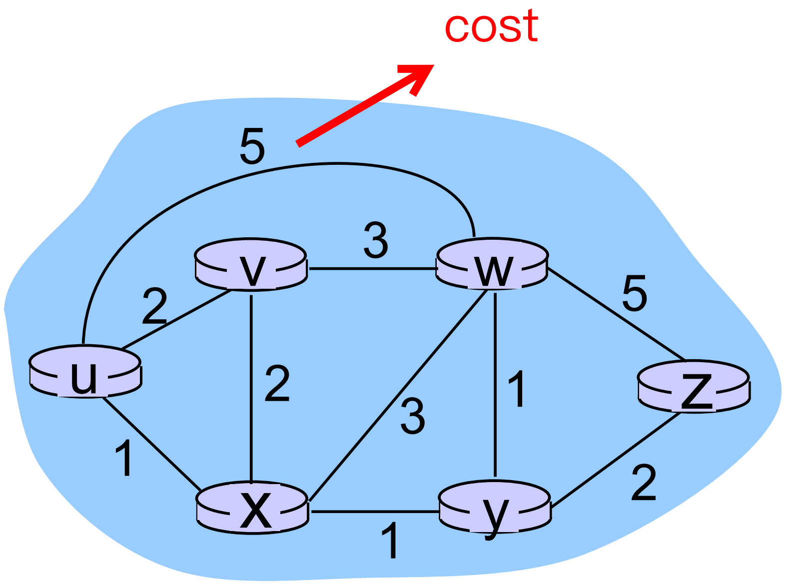

c(x,x') = cost of link (x,x')
e.g., c(w,z) = 5

cost of path $(x_1, x_2, x_3,\ldots, x_p) = c(x_1,x_2) + c(x_2,x_3) + \ldots + c(x_{p-1},x_p)$

key question: what is the least-cost path between u and z ?
routing algorithm: algorithm that finds that least cost path

Link Cost

- Typically simple: all links are equal
- Least-cost paths => shortest paths (hop count)
- Network operators add policy exceptions
  - Lower operational costs
  - Peering agreements
  - Security concerns

Some of the link maybe very expensive, like
those connected to satellite

Network layer, control plane: outline

5.1 introduction

5.2 routing protocols

- link state
- distance vector
- hierarchical routing

5.6 ICMP: The Internet Control Message Protocol

Routing algorithm classes

Link State (Global)

Simpler,
quickly

- Routers maintain cost of each link in the network

Limited size

- Connectivity/cost changes flooded to all routers
- Converges quickly (less inconsistency, looping, etc.)
- Limited network sizes

Can be scaled to very
large network

Distance Vector (Decentralised)

- Routers maintain next hop & cost of each destination.
- Connectivity/cost changes iteratively propagate form neighbour to neighbour
- Requires multiple rounds to converge
- Scales to large networks

Link State Routing

- Each node maintains its local "link state" (LS)
  - i.e., a list of its directly attached links and their costs

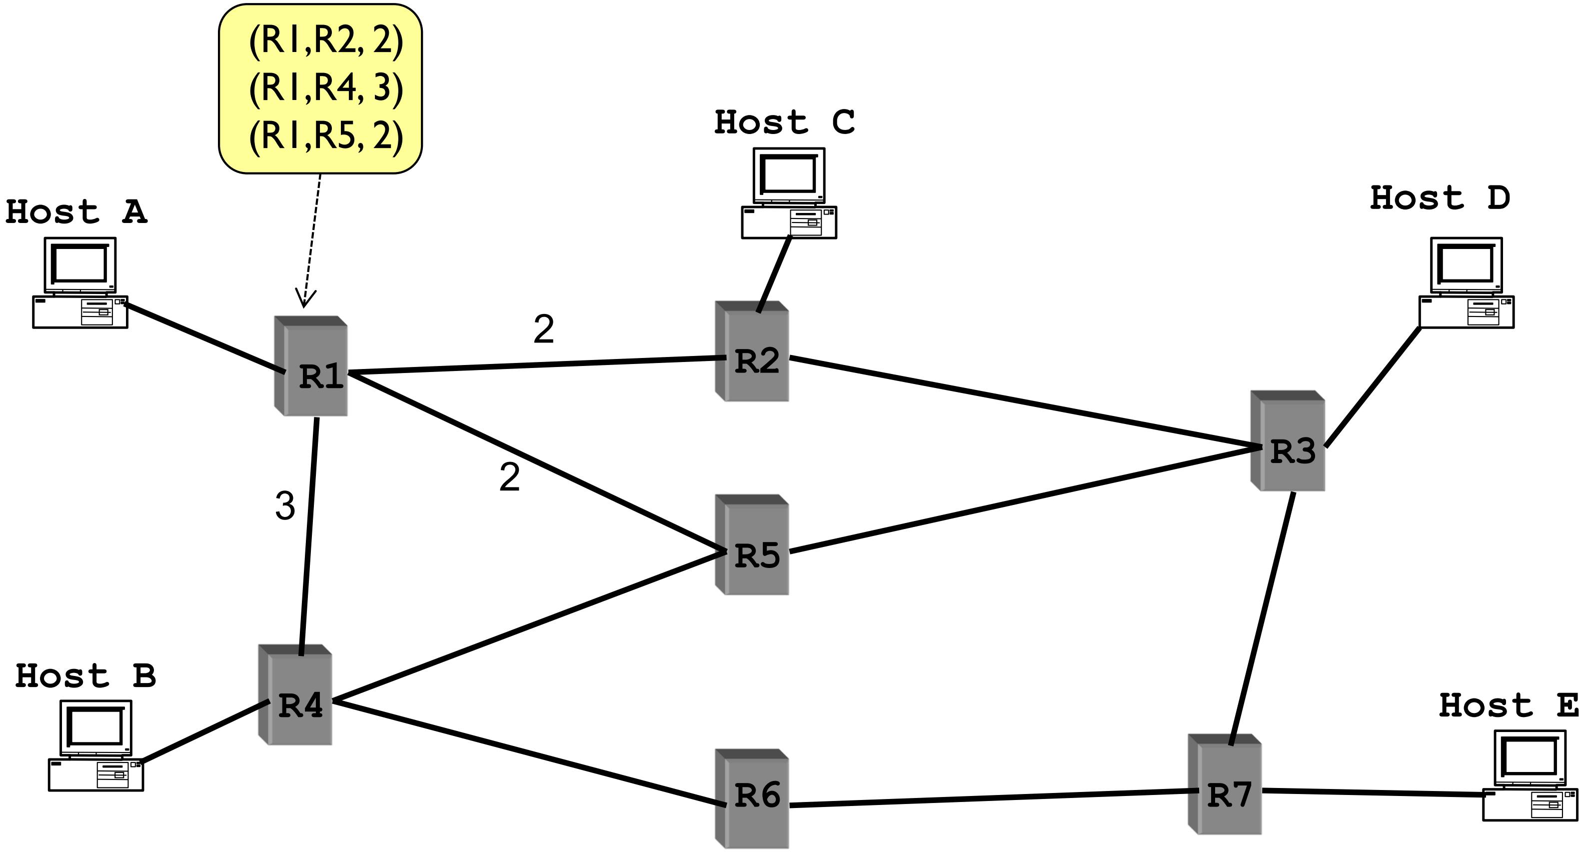
(R1,R2, 2)
(R1,R4, 3)
(R1,R5, 2)

Link State Routing

- Each node maintains its local "link state" (LS)
- Each node floods its local link state

Link state

  - on receiving a new LS message, a router forwards the message to all its neighbors other than the one it received the message from

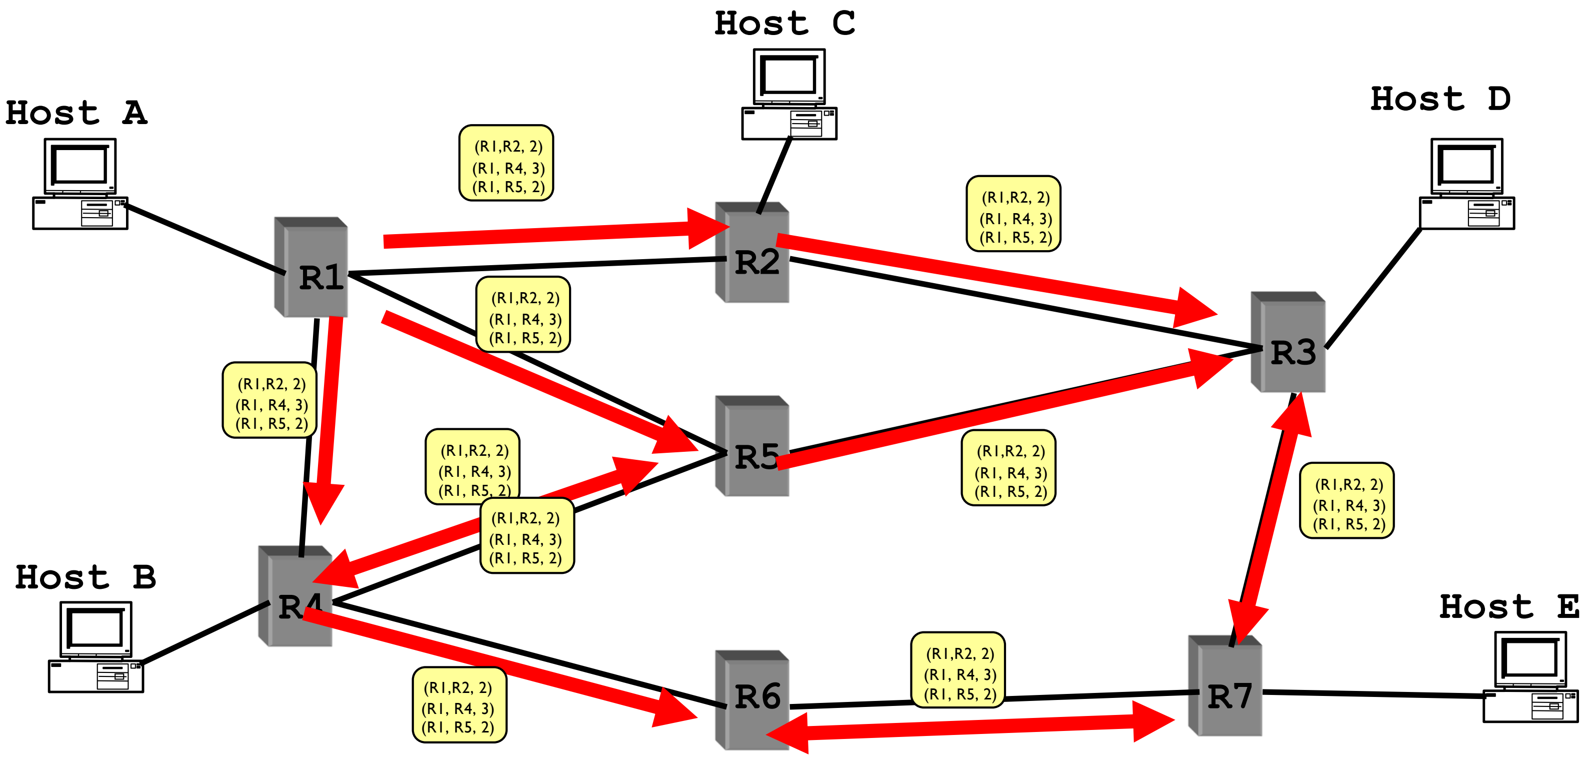

Flooding LSAs

- Routers transmit Link State Advertisement (LSA) on links
  - A neighbouring router forwards out on all links except incoming
  - Keep a copy locally; don't forward previously-seen LSAs
- Challenges
  - Packet loss
  - Out of order arrival
- Solutions
  - Acknowledgements and retransmissions
  - Sequence numbers
  - Time-to-live for each packet

Link State Routing

- Each node maintains its local "link state" (LS)
- Each node floods its local link state
- Eventually, each node learns the entire network topology
    - Can use Dijkstra's to compute the shortest paths between nodes

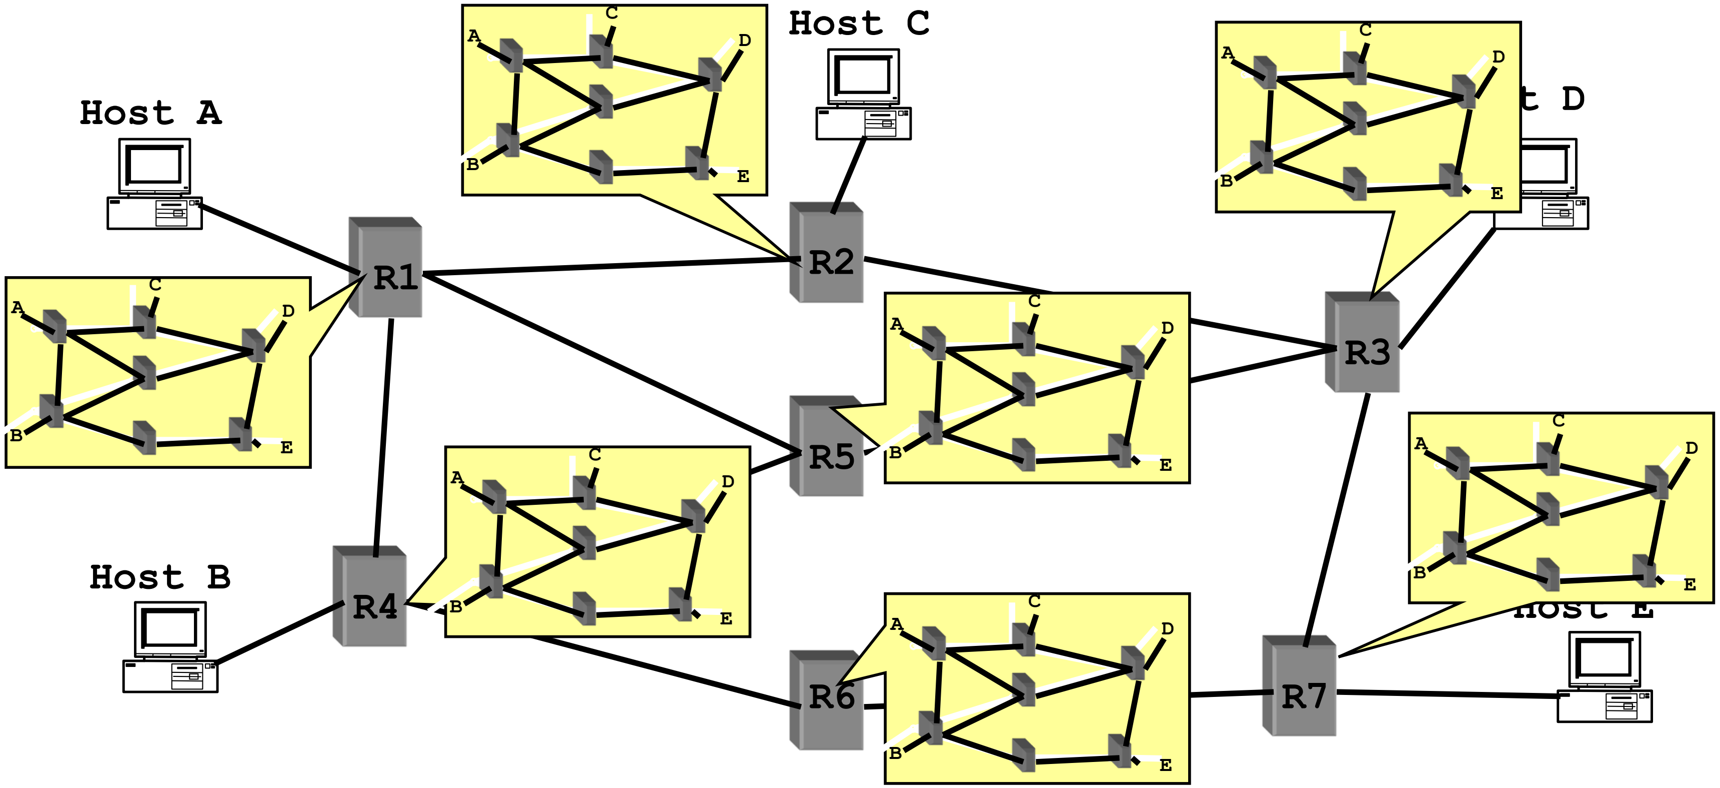

A Link-State Routing Algorithm

The number of iteration is the number of Node in the table

Dijkstra's algorithm

- net topology, link costs known to all nodes
  - accomplished via "link state broadcast"
  - all nodes have same info
- computes least cost paths from one node ('source") to all other nodes
  - gives *forwarding table* for that node
- iterative: after k iterations, know least cost path to k dest.'s

notation:

- c(x,y): link cost from node x to y; = ∞ if not direct neighbors
- D(v): current value of cost of path from source to dest. v
- p(v): predecessor node along path from source to v
- N': set of nodes whose least cost path definitively known

Dijsktra' s Algorithm

```
Initialization:
  N' = {u}
  for all nodes v
    if v adjacent to u
        then D(v) = c(u,v)
    else D(v) = ∞

Loop
  find w not in N' such that D(w) is a minimum
  add w to N'
  update D(v) for all v adjacent to w and not in N' :
      D(v) = min( D(v), D(w) + c(w,v) )
  /* new cost to v is either old cost to v or known
   shortest path cost to w plus cost from w to v */
until all nodes in N'
```

Example: Dijkstra's Algorithm

Step	Set N'	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1						
2						
3						
4						
5						

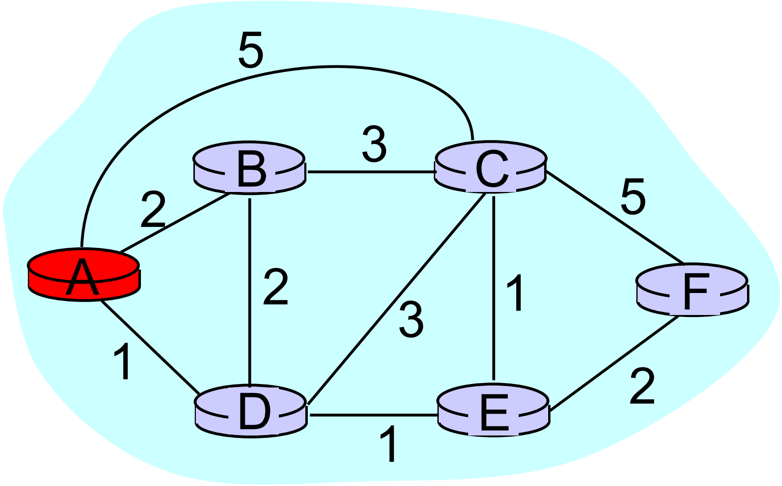

```
Initialization:
  N' = {A};
  for all nodes v
    if v adjacent to A
      then $D(v) = c(A,v)$;
      else $D(v) = \infty$;
…
```

Example: Dijkstra’s Algorithm

Step	Set N’	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1						
2						
3						
4						
5						

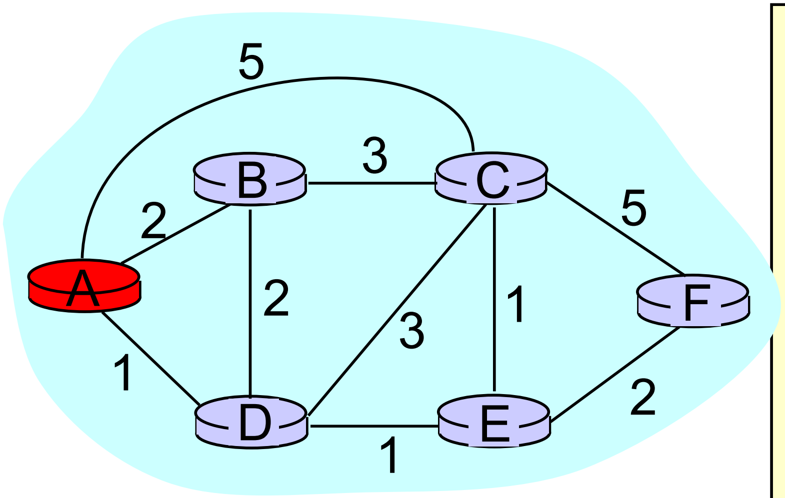

```
…
Loop
    find w not in N’ s.t. D(w) is a minimum;
    add w to N’;
    update D(v) for all v adjacent
        to w and not in N’:
    If D(w) + c(w,v) < D(v) then
        D(v) = D(w) + c(w,v); p(v) = w;
until all nodes in N’;
```

Example: Dijkstra’s Algorithm

Step	Set N’	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD					
2						
3						
4						
5						

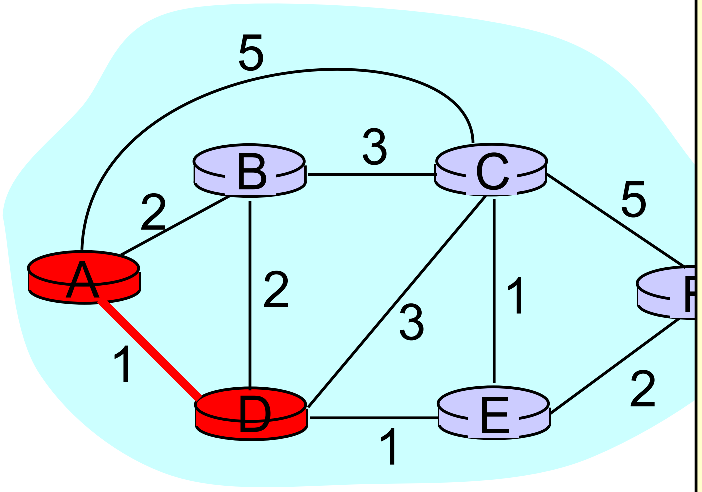

```
…
Loop
  find w not in N’ s.t. D(w) is a minimum;
  add w to N’;
  update D(v) for all v adjacent
    to w and not in N’:
  If D(w) + c(w,v) < D(v) then
     D(v) = D(w) + c(w,v); p(v) = w;
until all nodes in N’;
```

Example: Dijkstra’s Algorithm

Step	Set N’	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD	2, A	4,D		2,D	
2						
3						
4						
5						

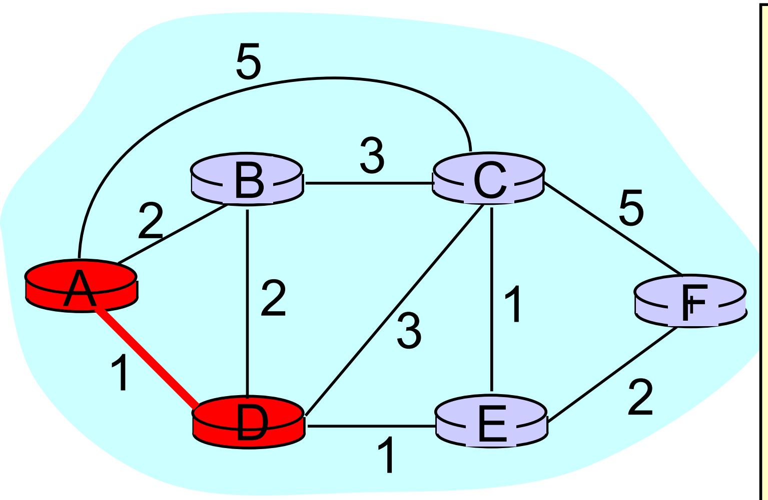

```
…
Loop
    find w not in N’ s.t. D(w) is a minimum;
    add w to N’;
    update D(v) for all v adjacent
      to w and not in N’:
    If D(w) + c(w,v) < D(v) then
        D(v) = D(w) + c(w,v); p(v) = w;
until all nodes in N’;
```

Example: Dijkstra's Algorithm

Step	Set N'	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD	2, A	4,D		2,D	
2	ADE	2, A	3,E			4,E
3						
4						
5						

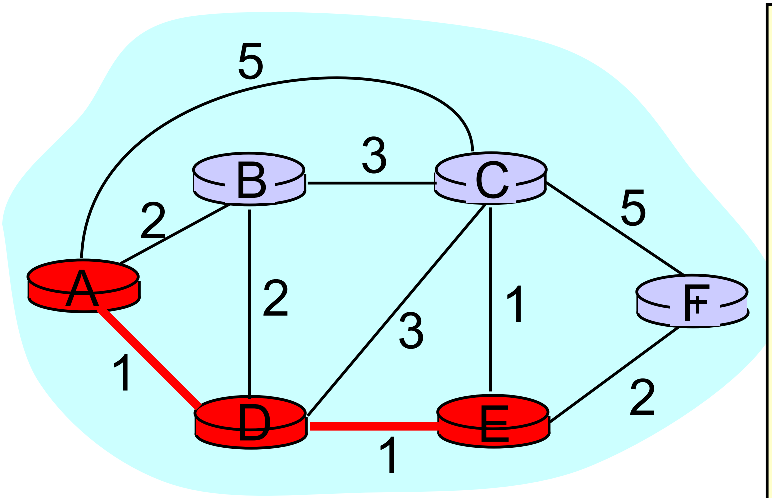

```
…
Loop
    find w not in N' s.t. D(w) is a minimum;
    add w to N';
   update D(v) for all v adjacent
      to w and not in N':
   If D(w) + c(w,v) < D(v) then
       D(v) = D(w) + c(w,v); p(v) = w;
    until all nodes in N';
```

Example: Dijkstra’s Algorithm

Step	Set N’	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD	2,A	4,D		2,D	
2	ADE	2,A	3,E			4,E
3	ADEB		3,E			4,E
4						
5						

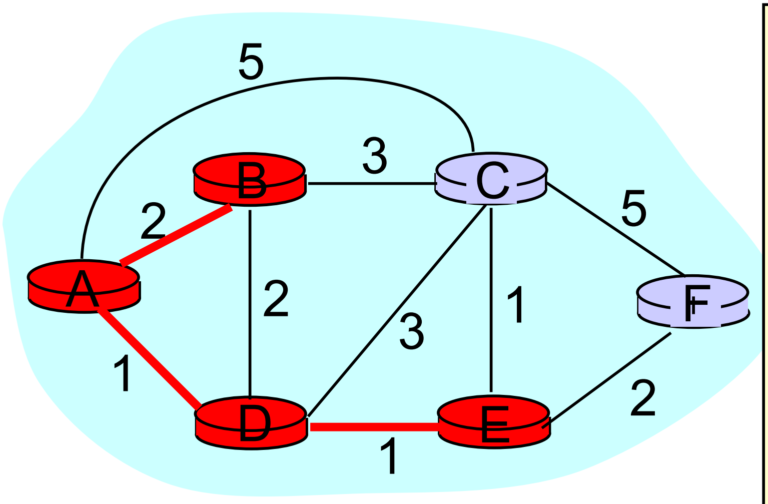

```
…
Loop
   find w not in N’ s.t. D(w) is a minimum;
   add w to N’;
  update D(v) for all v adjacent
     to w and not in N’:
  If D(w) + c(w,v) < D(v) then
     D(v) = D(w) + c(w,v); p(v) = w;
   until all nodes in N’;
```

Example: Dijkstra’s Algorithm

Step	Set N’	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD	2,A	4,D		2,D	
2	ADE	2,A	3,E			4,E
3	ADEB		3,E			4,E
4	ADEBC					4,E
5						

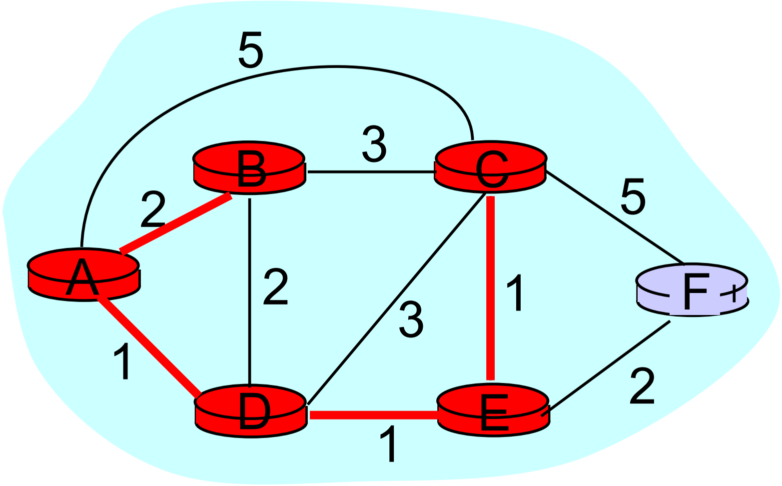

```
…
Loop
    find w not in N’ s.t. D(w) is a minimum;
    add w to N’;
    update D(v) for all v adjacent
        to w and not in N’:
    If D(w) + c(w,v) < D(v) then
        D(v) = D(w) + c(w,v); p(v) = w;
until all nodes in N’;
```

Example: Dijkstra’s Algorithm

Step	Set N’	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD	2,A	4,D		2,D	
2	ADE	2,A	3,E			4,E
3	ADEB		3,E			4,E
4	ADEBC					4,E
5	ADEBCF					

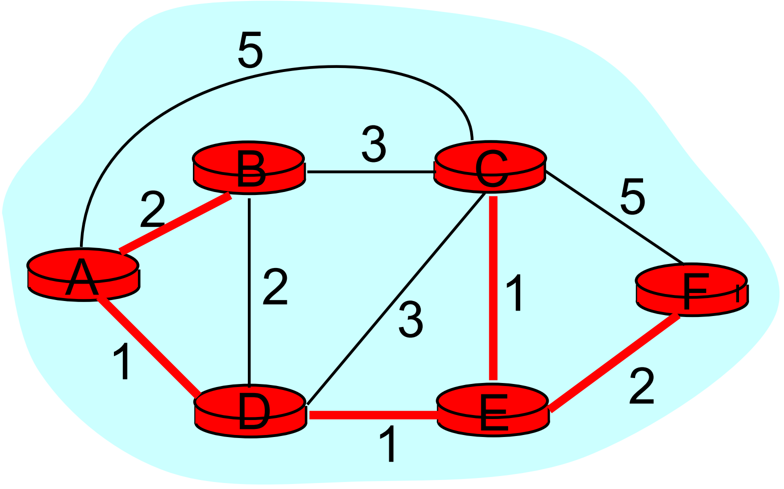

```
…
Loop
    find w not in N’ s.t. D(w) is a minimum;
    add w to N’;
    update D(v) for all v adjacent
      to w and not in N’:
    If D(w) + c(w,v) < D(v) then
        D(v) = D(w) + c(w,v); p(v) = w;
until all nodes in N’;
```

Example: Dijkstra’s Algorithm

Step	Set N’	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
0	A	2,A	5,A	1,A	∞	∞
1	AD		4,D		2,D	
2	ADE		3,E			4,E
3	ADEB					
4	ADEBC					
5	ADEBCF					

Once inclue all of the node, then we are done

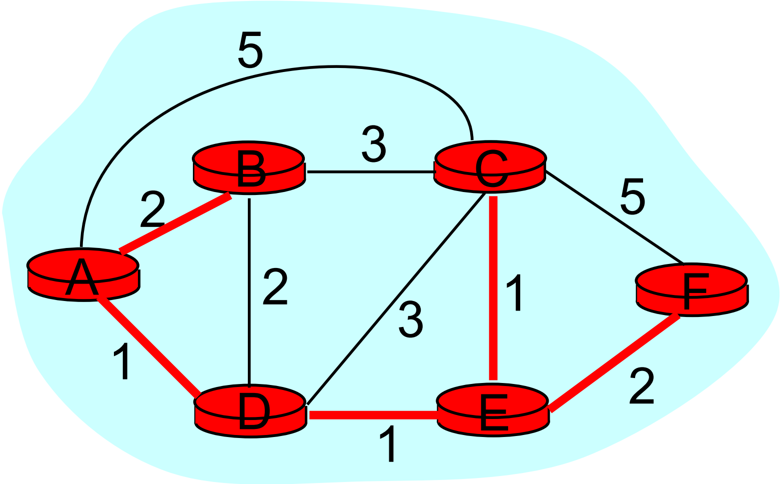

To determine path A → C (say), work backward from C via p(v)

The Forwarding Table

- Running Dijkstra at node A gives the shortest path from A to all destinations
- We then construct the *forwarding table*

resulting shortest-path tree from A:

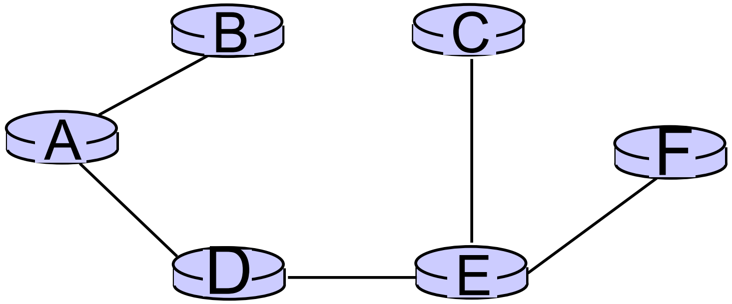

1. Flooding to learn the topology of the whole map
2. Get the shortest path to each node

Destination	Link
B	(A,B)
C	(A,D)
D	(A,D)
E	(A,D)
F	(A,D)

Issue #1: Scalability

- How many messages needed to flood link state messages?
  - O(N x E), where N is #nodes; E is #edges in graph
- Processing complexity for Dijkstra's algorithm?
  - O(N^2), because we check all nodes w not in N' at each iteration and we have O(N) iterations
- How many entries in the LS topology database? O(E)
- How many entries in the forwarding table? O(N)

Issue#2: Transient Disruptions

That is why we have TTL field

- Inconsistent link-state database
  - Some routers know about failure before others
  - The shortest paths are no longer consistent
  - Can cause **transient** forwarding loops

Some link may be corrupted unexpectively, then the network should do flooding again. But the flooding may take some time to reach all destination. So here comes the loop

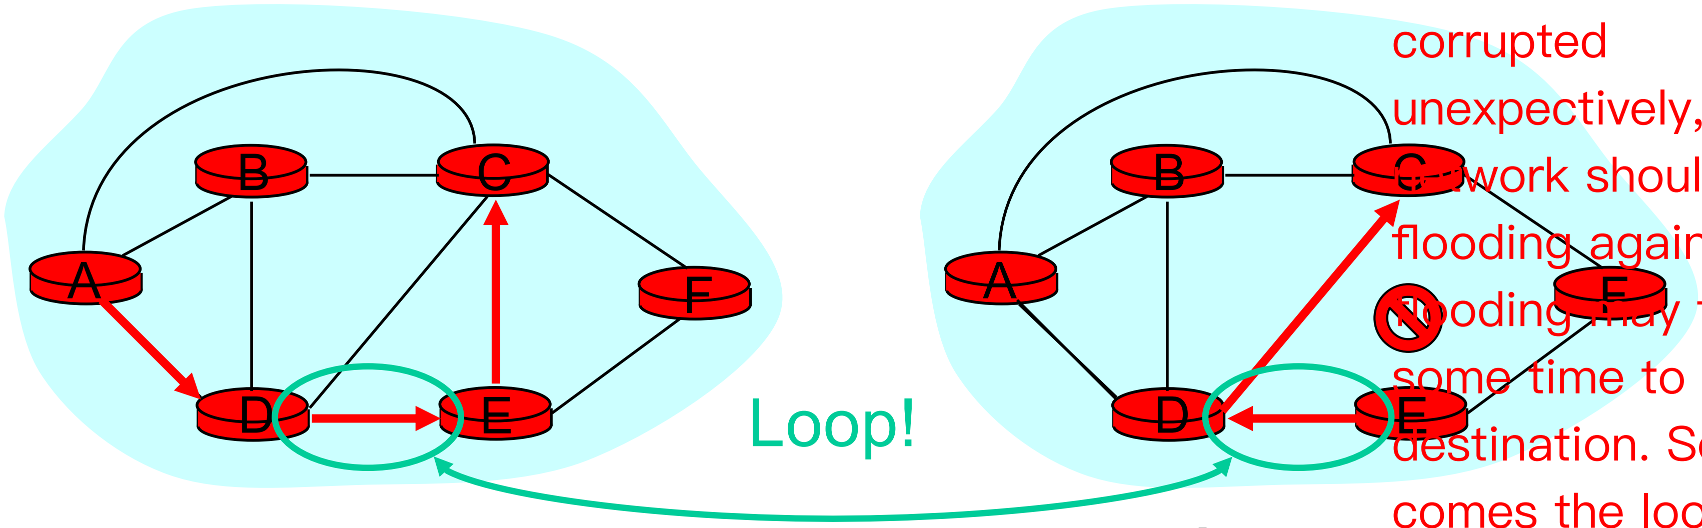

A and D think that this is the path to C

E thinks that this is the path to C

Oscillations

oscillations possible:

- e.g., suppose link cost equals amount of carried traffic:

as the traffic change frequently, the cost also change. As a result, the routers have to keep adjusting forwarding table

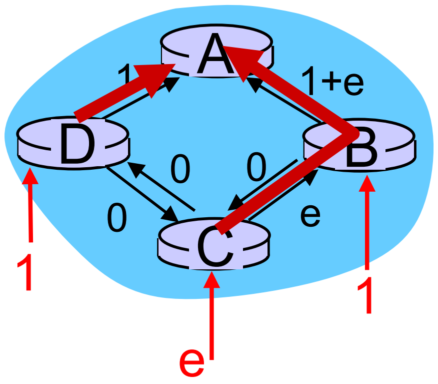

initially

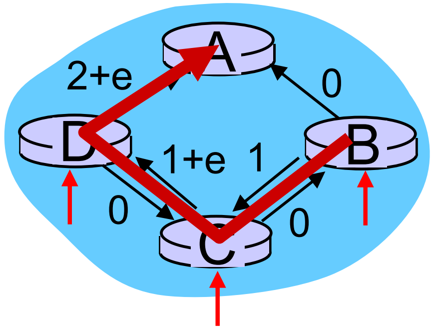

given these costs, find new routing…. resulting in new costs

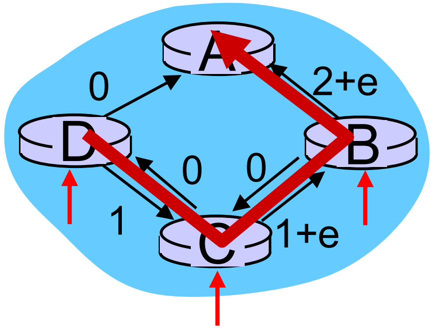

given these costs, find new routing…. resulting in new costs

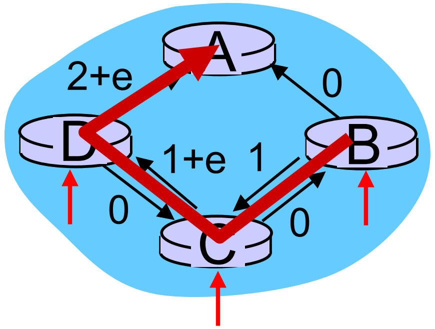

given these costs, find new routing…. resulting in new costs

That is a dynamic process

Network layer, control plane: outline

5.1 introduction

5.2 routing protocols

- link state
- distance vector
- hierarchical routing

5.6 ICMP: The Internet Control Message Protocol

Distance vector algorithm

Bellman-Ford equation

A kind of distributed protolco

let

$d_x(y)$:= cost of least-cost path from x to y

then

$$d_x(y) = \min_{v} \{c(x,v) + d_v(y) \}$$

cost from neighbor v to destination y

cost to neighbor v

min taken over all neighbors v of x

Bellman-Ford example

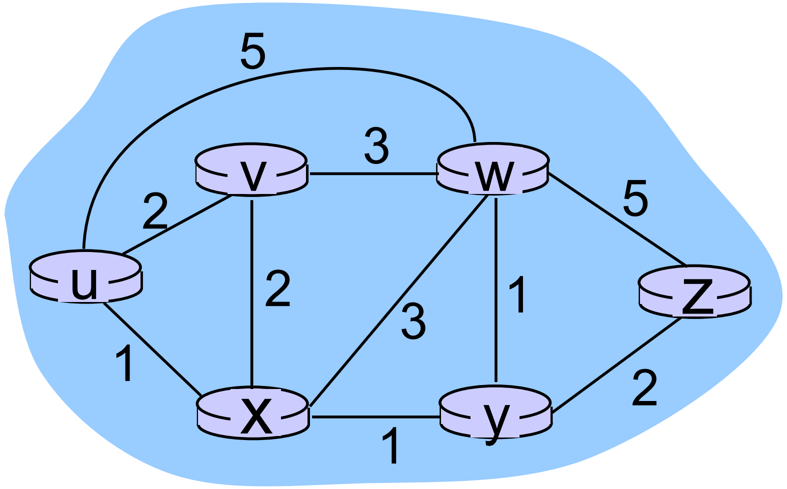

clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$\begin{aligned} d_u(z) = \min \{ & c(u,v) + d_v(z), \\ & c(u,x) + d_x(z), \\ & c(u,w) + d_w(z) \} \\ = \min \{ & 2 + 5, \\ & 1 + 3, \\ & 5 + 3\} = 4 \end{aligned}$$

node achieving minimum is next
hop in shortest path, used in forwarding table

How Distance-Vector (DV) works

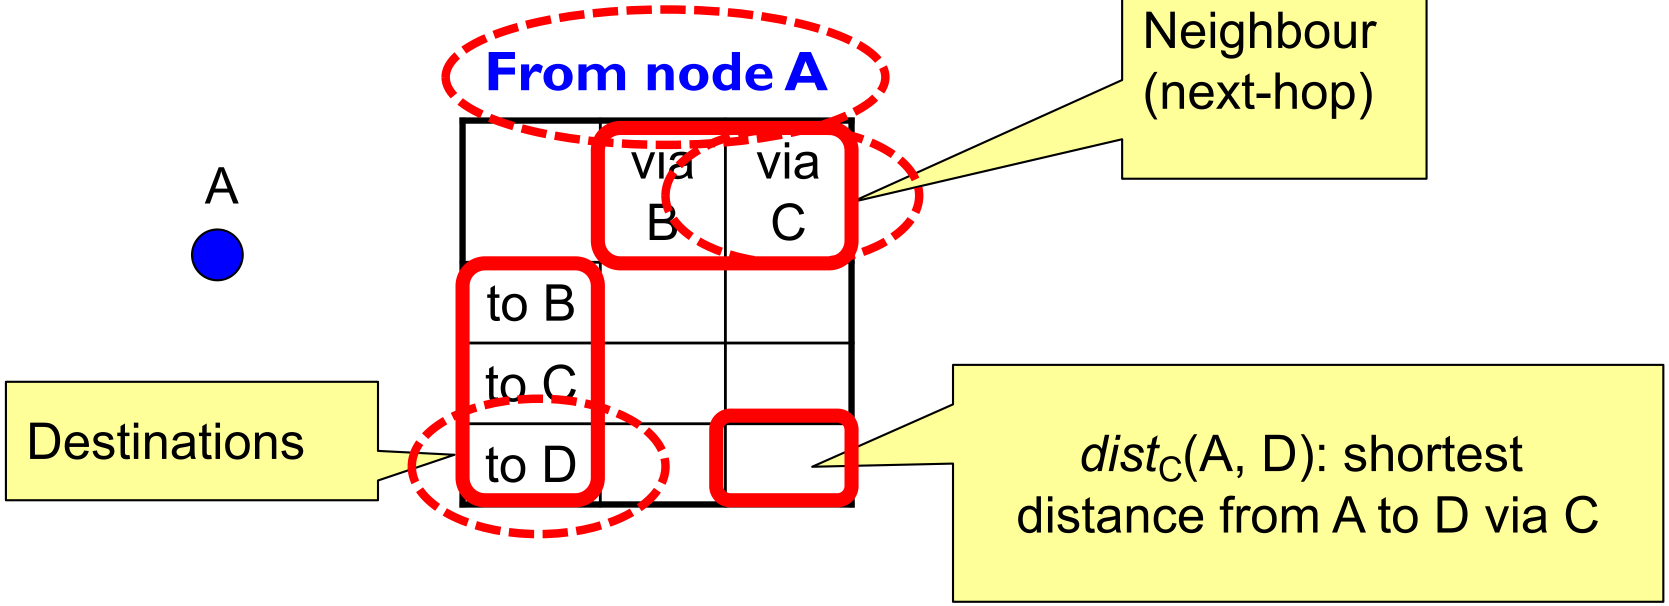

From node A

	via B	via C
to B		
to C		
to D		

Neighbour (next-hop)

Destinations

dist$_C$(A, D): shortest distance from A to D via C

Each router maintains its shortest distance to every destination via each of its neighbours

How Distance-Vector (DV) works

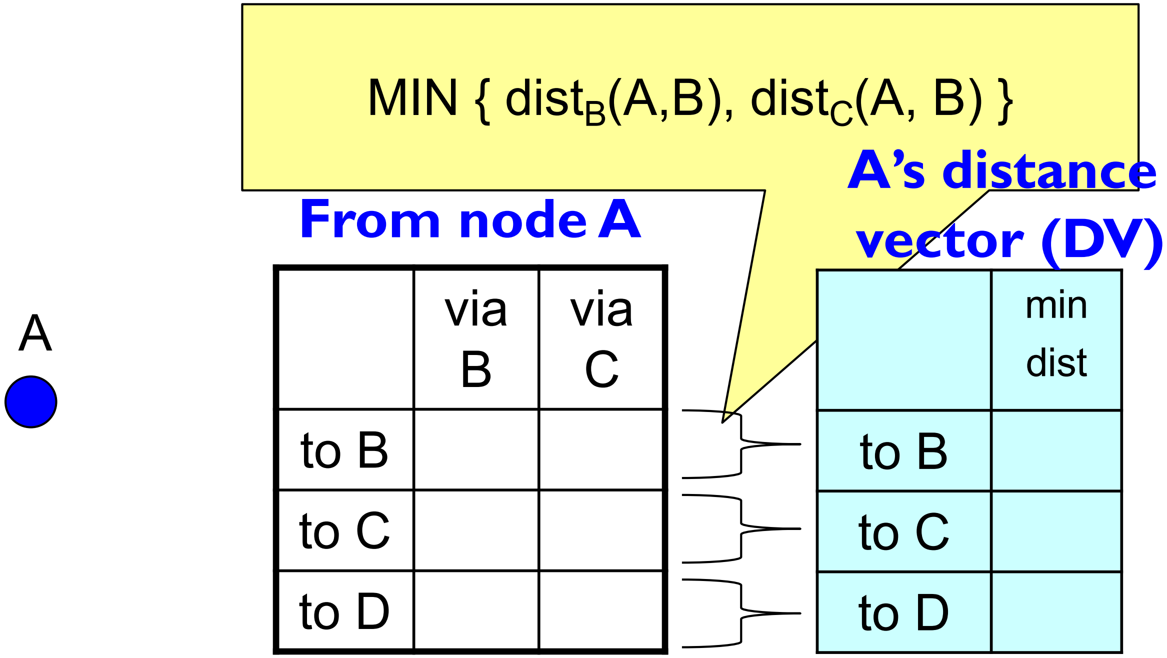

$\text{MIN} \{ \text{dist}_B(A,B), \text{dist}_C(A, B) \}$

From node A

	via B	via C
to B		
to C		
to D		

A's distance vector (DV)

	min dist
to B	
to C	
to D	

Each router computes its shortest distance to every destination via any of its neighbors

How Distance-Vector (DV) works

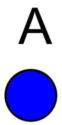

From node A

	via B	via C
to B	?	?
to C	?	?
to D	?	?

A’s DV

	min dist
to B	?
to C	?
to D	?

How does A initialize its dist() table and DV?

How Distance-Vector (DV) works

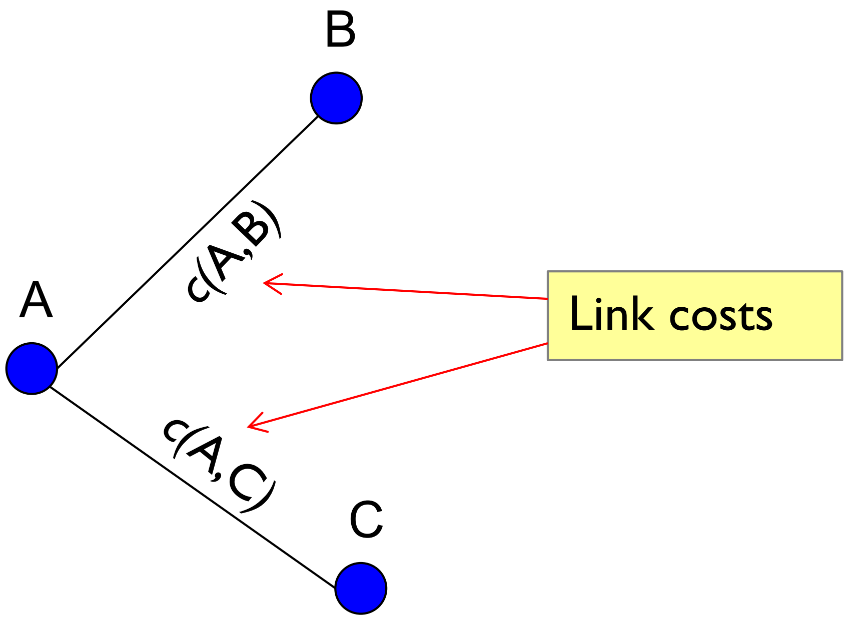

How does A initialize its dist() table and DV?

How Distance-Vector (DV) works

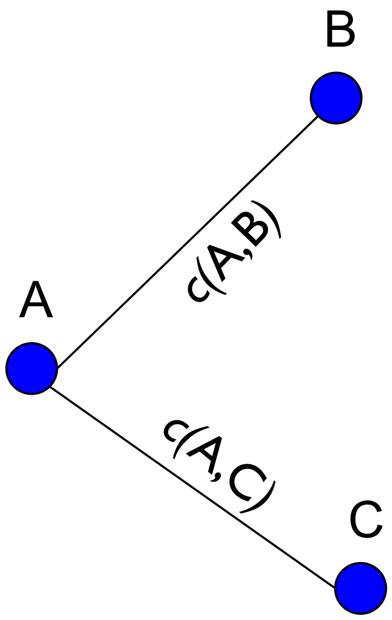

From node A

	via B	via C
to B	*c*(A,B)	∞
to C	∞	*c*(A,C)
to D	∞	∞

A's DV

	mindist
to B	c(A,B)
to C	c(A,C)
to D	∞

Exchange this table with its neighbors

Each router initializes its *dist*() table based on its immediate neighbors and link costs

How Distance-Vector (DV) works

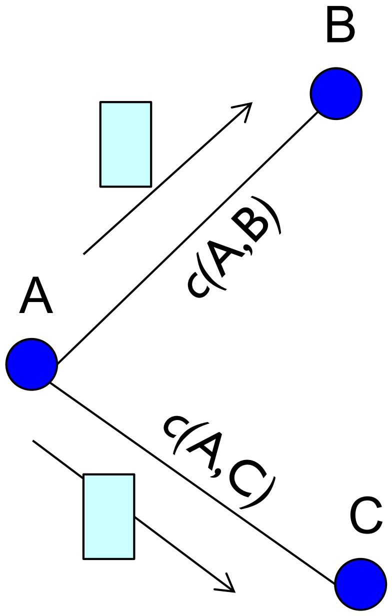

Assume that A's DV is as follows at some later time

From node A

	via B	via C
to B	*c*(A,B)	∞
to C	∞	*c*(A,C)
to D	∞	∞

A's DV

	mindist
to B	5
to C	6
to D	2

Each router sends its DV to its immediate neighbors

How Distance-Vector (DV) works

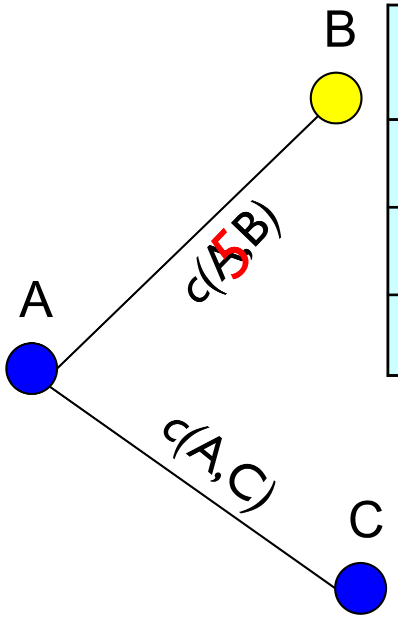

A’s DV

	min
to B	5
to C	6
to D	2

From node B

	via A	via C
to A	5	∞
to C	15	1
to D	∞	∞

B’s DV

	mindist
to A	5
to C	1
to D	∞

Routers process received DVs

How Distance-Vector (DV) works

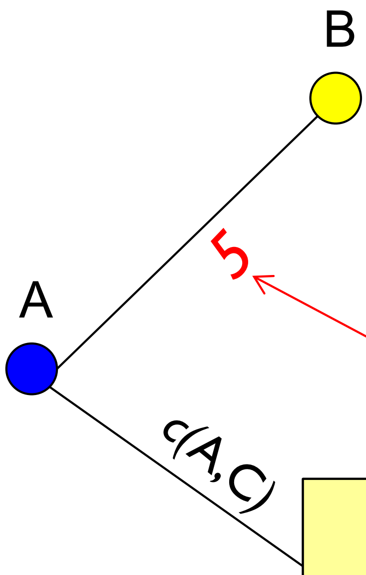

A's DV

	min
to B	5
to C	6
to D	2

From node B

	via A	via C
to A	5	∞
to C	11	1
to D	7	∞

B's DV

	mindist
to A	5
to C	1
to D	7

new = c(B,A) + *mindist*(A, D)

Routers process received DVs

And repeat…

Distance Vector Routing

- Each router knows the links to its neighbors
- Each router has provisional "shortest path" to every other router -- its distance vector (DV)
- Routers exchange this DV with their neighbors

Finally they would converge, which means the router who have the full network of destinations in the table. And they has the shortest path. In a way, the packets would be always guided to the shortest path

- Routers look over the set of options offered by their neighbors and select the best one

Y

- Iterative process converges to set of shortest paths

Distance vector routing

iterative, asynchronous:
each local iteration caused by:

- local link cost change
- DV update message from neighbor

distributed:

- each node notifies neighbors *only* when its DV changes
  - neighbors then notify their neighbors if necessary

each node:

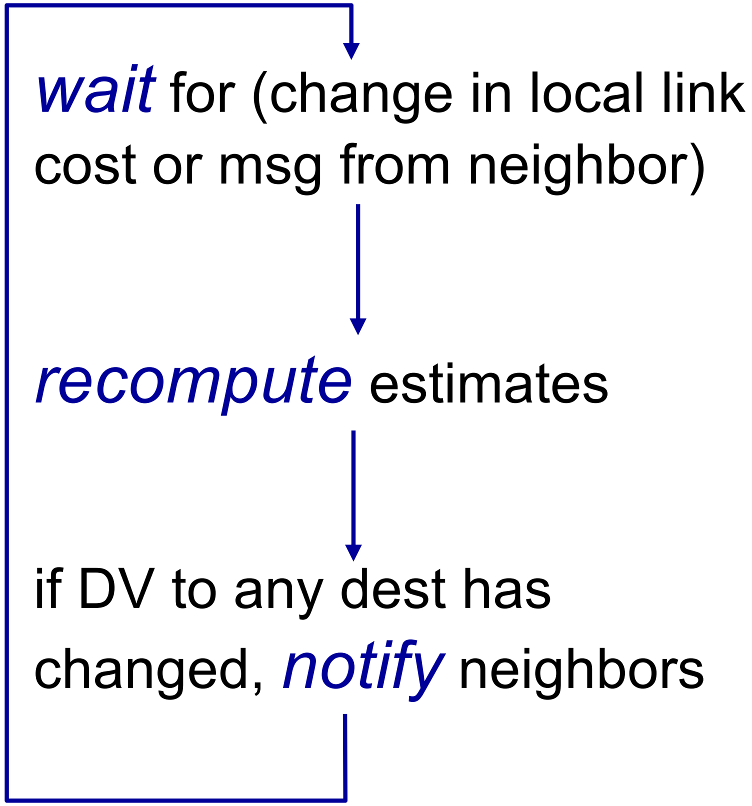

wait for (change in local link cost or msg from neighbor)

recompute estimates

if DV to any dest has changed, *notify* neighbors

Distance Vector

- c(i,j): link cost from node *i* to *j*
- $dist_Z(A,V)$: shortest dist. from A to V via Z
- mindist(A,V): shortest dist. from A to V

```
At node A
Initialization:
  for all destinations V do
      if V is neighbor of A
          $dist_V(A, V)$ = mindist(A,V) = c(A,V);
      else
          $dist_V(A, V)$ = mindist(A,V) = $\infty$;
  send mindist(A, *) to all neighbors
loop:
  wait (until A sees a link cost change to neighbor V  /* case 1 */
        or until A receives mindist(V,*) from neighbor V)   /* case 2 */
  if (c(A,V) changes by $\pm d$)  /* $\Leftarrow$ case 1 */
      for all destinations Y do
          $dist_V(A,Y)$ = $dist_V(A,Y) \pm d$
  else /* $\Leftarrow$ case 2: */
      for all destinations Y do
          $dist_V(A,Y)$ = c(A,V) + mindist(V, Y);
  update mindist(A,*)
  if (there is a change in mindist(A, *))
      send mindist(A, *) to all neighbors
forever
```

Distance Vector

- c(i,j): link cost from node *i* to *j*
- $dist_Z(A,V)$: shortest dist. from A to V via Z
- mindist(A,V): shortest dist. from A to V

```
At node A
Initialization:
  for all destinations V do
     if V is neighbor of A
        $dist_V(A, V)$ = mindist(A,V) = c(A,V);
     else
        $dist_V(A, V)$ = mindist(A,V) = $\infty$;
   send mindist(A, *) to all neighbors
loop:
  wait (until A sees a link cost change to neighbor V  /* case 1 */
       or until A receives mindist(V,*) from neighbor V)   /* case 2 */
   if (c(A,V) changes by $\pm d$)  /* $\Leftarrow$ case 1 */
     for all destinations Y do
           $dist_V(A,Y)$ =  $dist_V(A,Y) \pm d$
   else /* $\Leftarrow$ case 2: */
     for all destinations Y do
           $dist_V(A,Y)$ = c(A,V) + mindist(V, Y);
   update mindist(A,*)
  if (there is a change in mindist(A, *))
       send mindist(A,  *) to all neighbors
  forever
```

Example: Initialization

from Node B

	via A	via C	via D	min dist
to A	2	∞	∞	2
to B	-	-	-	0
to C	∞	1	∞	1
to D	∞	∞	3	3

from Node D

	via B	via C	min dist
to A	∞	∞	∞
to B	3	∞	3
to C	∞	1	1
to D	-	-	0

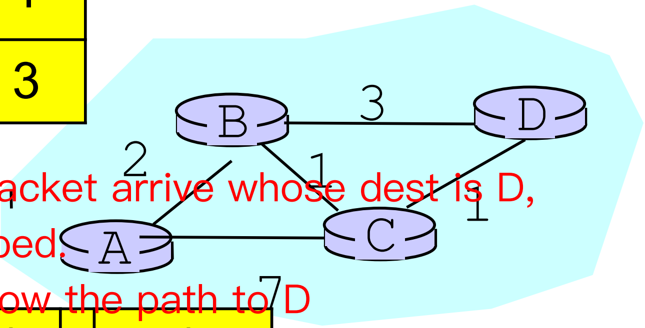

At this moment, if a packet arrive whose dest is D, then it would be dropped. Because A doesn't know the path to D

from Node A

	via B	via C	min dist	min dist
to A	-	-	0	0
to B	2	∞	2	2
to C	∞	7	7	7
to D	∞	∞	∞	∞

from Node C

	via A	via B	via D	min dist
to A	7	∞	∞	7
to B	∞	1	∞	1
to C	-	-	-	0
to D	∞	∞	1	1

Example: C sends update to A

from Node B

	via A	via C	via D	min dist
to A	2	∞	∞	2
to B	-	-	-	0
to C	∞	1	∞	1
to D	∞	∞	3	3

from Node D

	via B	via C	min dist
to A	∞	∞	∞
to B	3	∞	3
to C	∞	1	1
to D	-	-	0

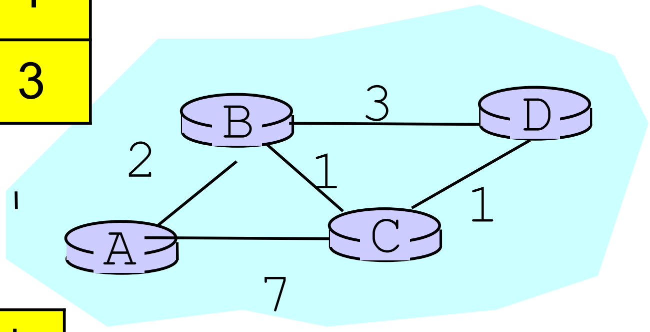

from Node A

	via B	via C	min dist
to A	-	-	0
to B	2	∞	2
to C	∞	7	7
to D	∞	∞	∞

from Node C

	via A	via B	via D	min dist
to A	7	∞	∞	7
to B	∞	1	∞	1
to C	-	-	-	0
to D	∞	∞	1	1

Example: C sends update to A

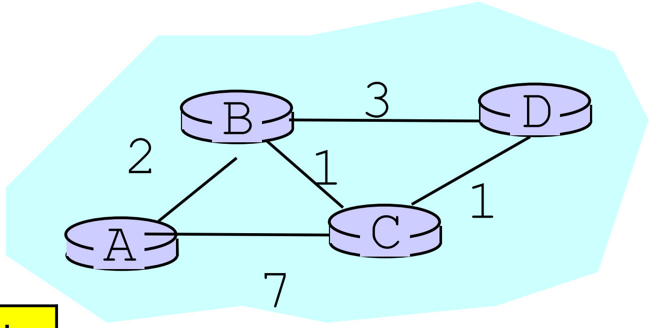

from Node A

	via B	via C
to A	-	-
to B	2	∞
to C	∞	7
to D	∞	∞

min dist
0
2
7
∞

min dist
7
1
0
1

Example: C sends update to A

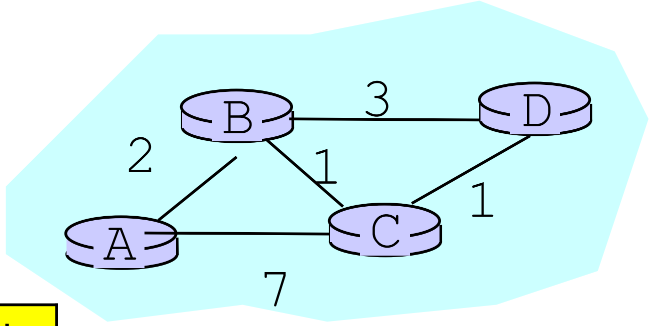

from Node A

	via B	via C
to A	-	-
to B	2	**8**
to C	∞	7
to D	∞	**8**

min dist
0
2
7
∞

min dist
7
1
0
1

Example: C sends update to A

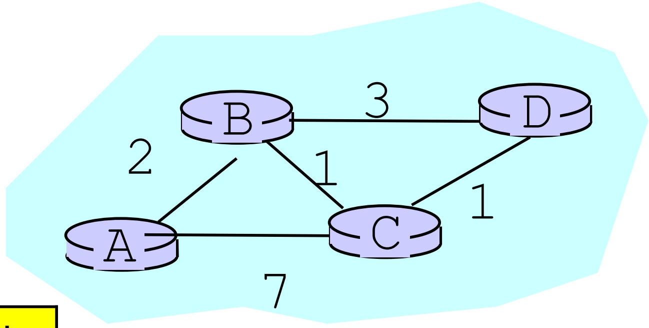

from Node A

	via B	via C
to A	-	-
to B	2	**8**
to C	∞	7
to D	∞	**8**

min dist
0
2
7
8

Example: now B sends update to A

from Node B

	via A	via C	via D
to A	2	∞	∞
to B	-	-	-
to C	∞	1	∞
to D	∞	∞	3

min dist
2
0
1
3

from Node D

	via B	via C
to A	∞	∞
to B	3	∞
to C	∞	1
to D	-	-

min dist
∞
3
1
0

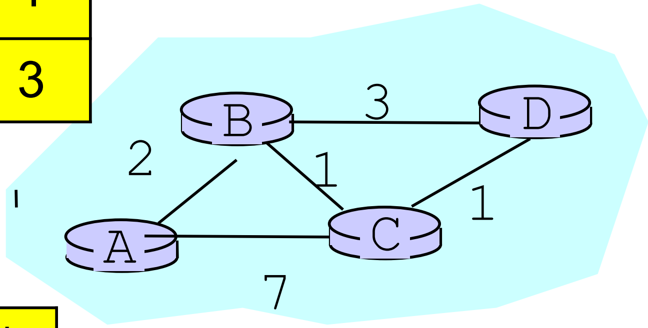

from Node A

	via B	via C
to A	-	-
to B	2	8
to C	∞	7
to D	∞	8

min dist
0
2
7
8

from Node C

	via A	via B	via D
to A	7	∞	∞
to B	∞	1	∞
to C	-	-	-
to D	∞	∞	1

min dist
7
1
0
1

Example: now B sends update to A

from Node B

	via A	via C	via D	min dist
to A	2	∞	∞	2
to B	-	-	-	0
to C	∞	1	∞	1
to D	∞	∞	3	3

from Node D

	via B	via C	min dist
to A	∞	∞	∞
to B	3	∞	3
to C	∞	1	1
to D	-	-	0

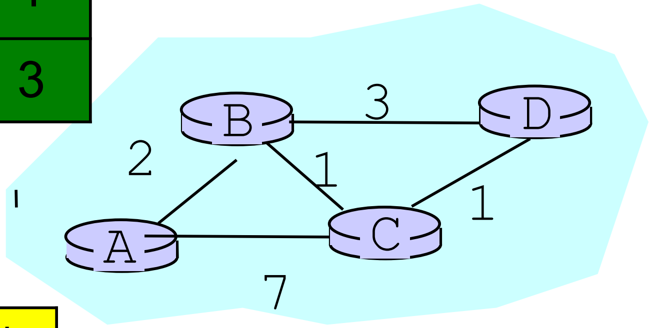

from Node A

	via B	via C	min dist
to A	-	-	0
to B	2	8	2
to C	∞	7	7
to D	∞	8	8

from Node C

	via A	via B	via D	min dist
to A	7	∞	∞	7
to B	∞	1	∞	1
to C	-	-	-	0
to D	∞	∞	1	1

Example: now B sends update to A

from Node B

	via A	via C	via D	min dist
to A	2	∞	∞	2
to B	-	-	-	0
to C	∞	1	∞	1
to D	∞	∞	3	3

from Node D

	via B	via C	min dist
to A	∞	∞	∞
to B	3	∞	3
to C	∞	1	1
to D	-	-	0

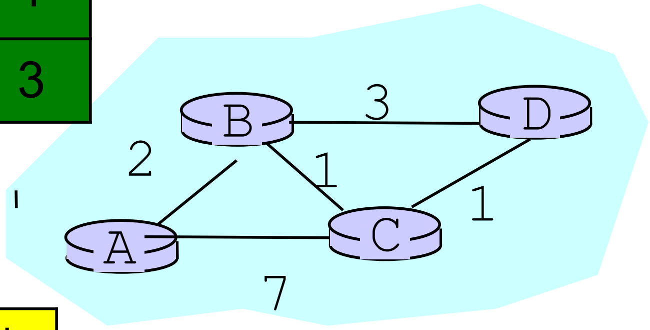

from Node A

	via B	via C	min dist
to A	-	-	0
to B	2	8	2
to C	3	7	7
to D	5	8	8

Make sure you know why this is 5, not 4!

from Node C

	via	via	via	min dist
				7
				1
to C	-	-	-	0
to D	∞	∞	1	1

Example: now B sends update to A

from Node B

	via A	via C	via D	min dist
to A	2	∞	∞	2
to B	-	-	-	0
to C	∞	1	∞	1
to D	∞	∞	3	3

from Node D

	via B	via C	min dist
to A	∞	∞	∞
to B	3	∞	3
to C	∞	1	1
to D	-	-	0

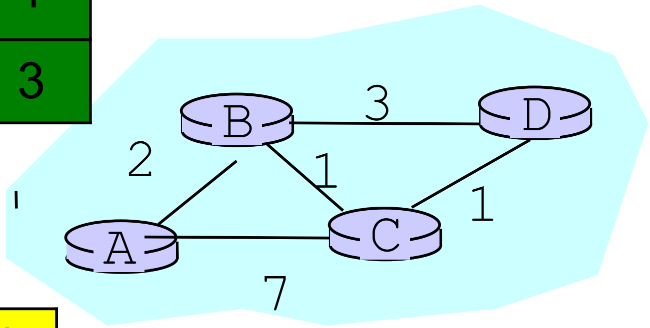

from Node A

	via B	via C	min dist
to A	-	-	0
to B	2	8	2
to C	**3**	7	**3**
to D	**5**	8	**5**

from Node C

	via A	via B	via D	min dist
to A	7	∞	∞	7
to B	∞	1	∞	1
to C	-	-	-	0
to D	∞	∞	1	1

All nodes know the best two-hop paths.

Make sure you believe this

from Node B

	via A	via C	via D	min dist
to A	2	8	∞	2
to B	-	-	-	0
to C	9	1	4	1
to D	∞	2	3	**2**

from Node D

	via B	via C	min dist
to A	5	8	**5**
to B	3	2	**2**
to C	4	1	1
to D	-	-	0

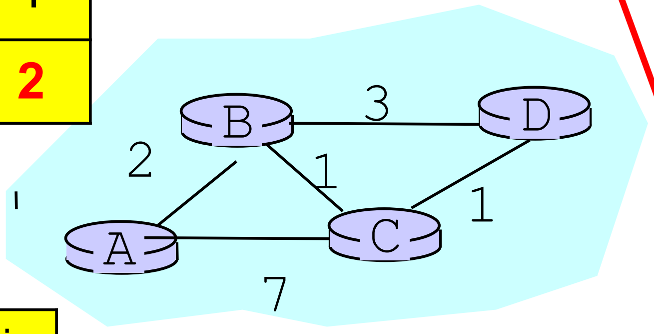

from Node A

	via B	via C	min dist
to A	-	-	0
to B	2	8	2
to C	3	7	**3**
to D	5	8	**5**

from Node C

	via A	via B	via D	min dist
to A	7	3	∞	**3**
to B	9	1	4	1
to C	-	-	-	0
to D	∞	4	1	1

Example: Now A sends update to B

from Node B

	via A	via C	via D	min dist
to A	2	8	∞	2
to B	-	-	-	0
to C	9	1	4	1
to D	∞	2	3	2

from Node D

	via B	via C	min dist
to A	5	8	5
to B	3	2	2
to C	4	1	1
to D	-	-	0

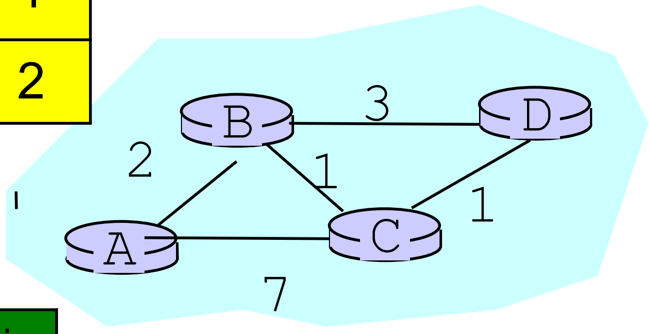

from Node A

	via B	via C	min dist
to A	-	-	0
to B	2	8	2
to C	3	7	3
to D	5	8	5

from Node C

	via A	via B	via D	min dist
to A	7	3	∞	3
to B	9	1	4	1
to C	-	-	-	0
to D	∞	4	1	1

Example: Now

Updated

from Node B

	via A	via C	via D
to A	2	8	∞
to B	-	-	
to C	5	1	4
to D	7	2	3

mi d
0
1
2

from Node D

	via B	via C
to A	5	8
to B	3	2
to C	4	1
to D	-	-

min dist
5
2
1
0

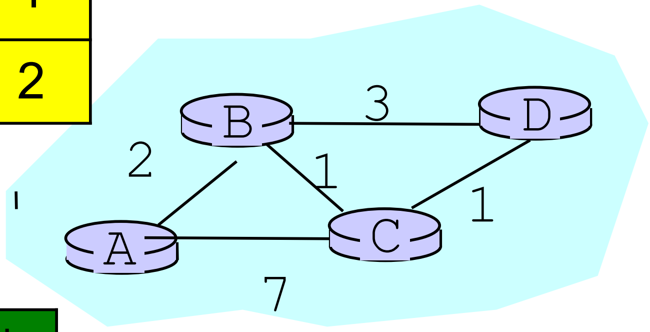

from Node A

	via B	via C
to A	-	-
to B	2	8
to C	3	7
to D	5	8

min dist
0
2
3
5

from Node C

	via A	via B	via D
to A	7	3	∞
to B	9	1	4
to C	-	-	-
to D	∞	4	1

min dist
3
1
0
1

Check: *All nodes know the best three-hop paths.*

from Node B

	via A	via C	via D	min dist
to A	2	4	8	2
to B	-	-	-	0
to C	5	1	4	1
to D	7	2	3	2

from Node D

	via B	via C	min dist
to A	5	4	4
to B	3	2	2
to C	4	1	1
to D	-	-	0

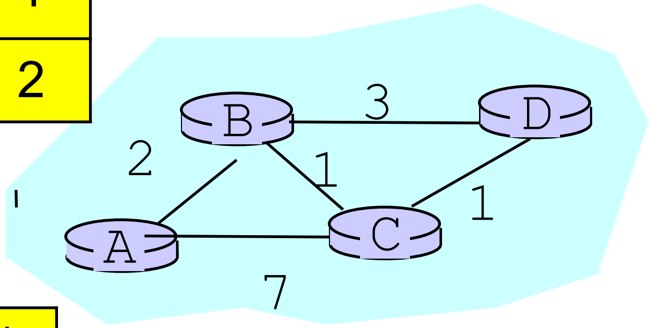

from Node A

	via B	via C	min dist
to A	-	-	0
to B	2	8	2
to C	3	7	3
to D	4	8	4

from Node C

	via A	via B	via D	min dist
to A	7	3	6	3
to B	9	1	3	1
to C	-	-	-	0
to D	12	3	1	1

Check

Example: End of 3rd Full Exchange

No further change in DVs → Convergence!

from Node B

	via A	via C	via D	min dist
to A	2	4	7	2
to B	-	-	-	0
to C	5	1	4	1
to D	6	2	3	2

from Node D

	via B	via C	min dist
to A	5	4	4
to B	3	2	2
to C	4	1	1
to D	-	-	0

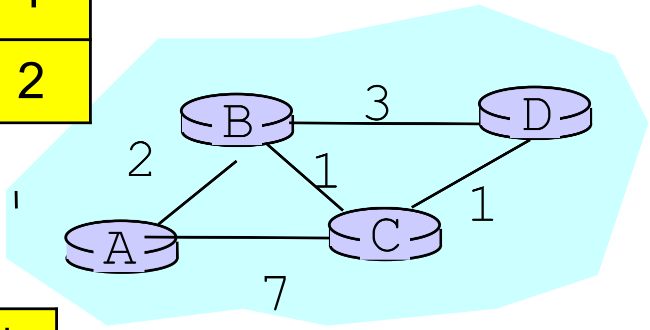

from Node A

	via B	via C	min dist
to A	-	-	0
to B	2	8	2
to C	3	7	3
to D	4	8	4

from Node C

	via A	via B	via D	min dist
to A	7	3	5	3
to B	9	1	3	1
to C	-	-	-	0
to D	11	3	1	1

Intuition

- Initial state: best one-hop paths
- One simultaneous round: best two-hop paths
- Two simultaneous rounds: best three-hop paths
- …
- Kth simultaneous round: best (k+1) hop paths

- Must eventually converge
  - as soon as it reaches longest best path
- …..but how does it respond to changes in cost?

Problems with Distance Vector

- A number of problems can occur in a network using distance vector algorithm
- Most of these problems are caused by slow convergence or routers converging on incorrect information
- *Convergence* is the time during which all routers come to an agreement about the best paths through the internetwork
  - whenever topology changes there is a period of instability in the network as the routers converge
- Reacts rapidly to good news, but leisurely to bad news

DV: Link Cost Changes

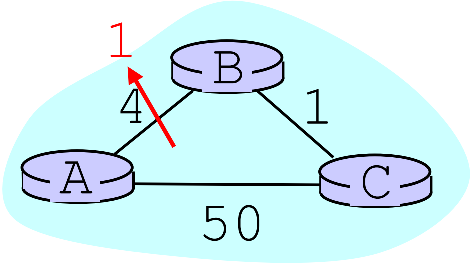

via

to

Node A

	Stable state		A-B changed		A sends its DV to B, C		B sends its DV to A, C		C sends its DV to A, B	
	B	C	B	C	B	C	B	C	B	C
B	4	51	1	51	1	51	1	51	1	51
C	5	50	2	5						50

Node B

	Stable state		A-B changed		A sends its DV to B, C		B sends its DV to A, C		C sends its DV to A, B	
	A	C								C
A	4	6	1	6						3
C	6	1	6	1	3	1	3	1	3	1

Node C

	Stable state		A-B changed		A sends its DV to B, C		B sends its DV to A, C		C sends its DV to A, B	
	A	B	A	B	A	B	A	B	A	B
A	50	5	50	5	50	5	50	2	50	2
B	54	1	54	1	51	1	51	1	51	1

deduct 3 from distances $dist_B(A,*)$ and $dist_A(B,*)$

Link cost changes here

"good news travels fast"

DV: Link Cost Changes

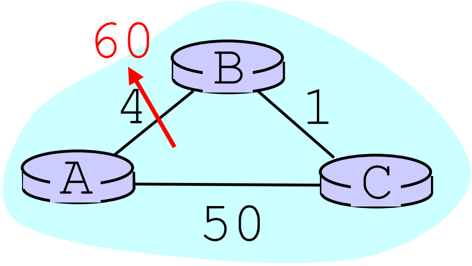

Stable state

A-B changed

via

to

Node A

Stable state:

	B	C
B	4	51
C	5	50

A-B changed:

	B	C
B	60	51
C	61	50

Node B

Stable state:

	A	C
A	4	6
C	6	1

A-B changed:

	A	C
A	60	6
C	65	1

Node C

Stable state:

	A	B
A	50	5
B	54	1

A-B changed:

	A	B
A	50	5
B	54	1

add 56 to distances
$\text{dist}_B(A,*)$ and $\text{dist}_A(B,*)$

Link cost changes here

DV: Link Cost Changes

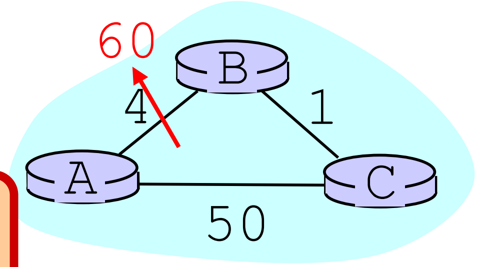

This is the “Counting to Infinity” Problem

Stable state

via

to

Node A

	B	C
B	4	51
C	5	50

Node B

	A	C
A	4	6
C	9	1

Node C

	A	B
A	50	5
B	54	1

A-B changed

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	6
C	65	1

Node C

	A	B
A	50	5
B	54	1

A sends its DV to B, C

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	6
C	110	1

Node C

	A	B
A	50	5
B	101	1

B sends its DV to A, C

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	6
C	110	1

Node C

	A	B
A	50	7
B	101	1

C sends its DV to A, B

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	8
C	110	1

Node C

	A	B
A	50	7
B	101	1

Link cost changes here

“bad news travels slowly”
(not yet converged)

The “Poisoned Reverse” Rule

- Heuristic to avoid count-to-infinity

- If B routes via C to get to A:
  - B tells C its (B’s) distance to A is infinite (so C won’t route to A via B)

DV: Poisoned Reverse

If B routes through C to get to A:
B tells C its (B's) distance to A is infinite

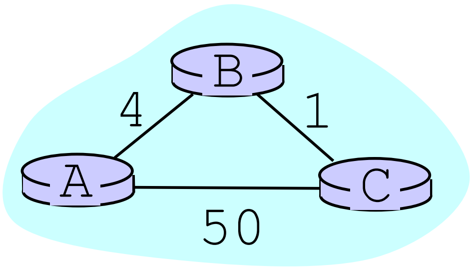

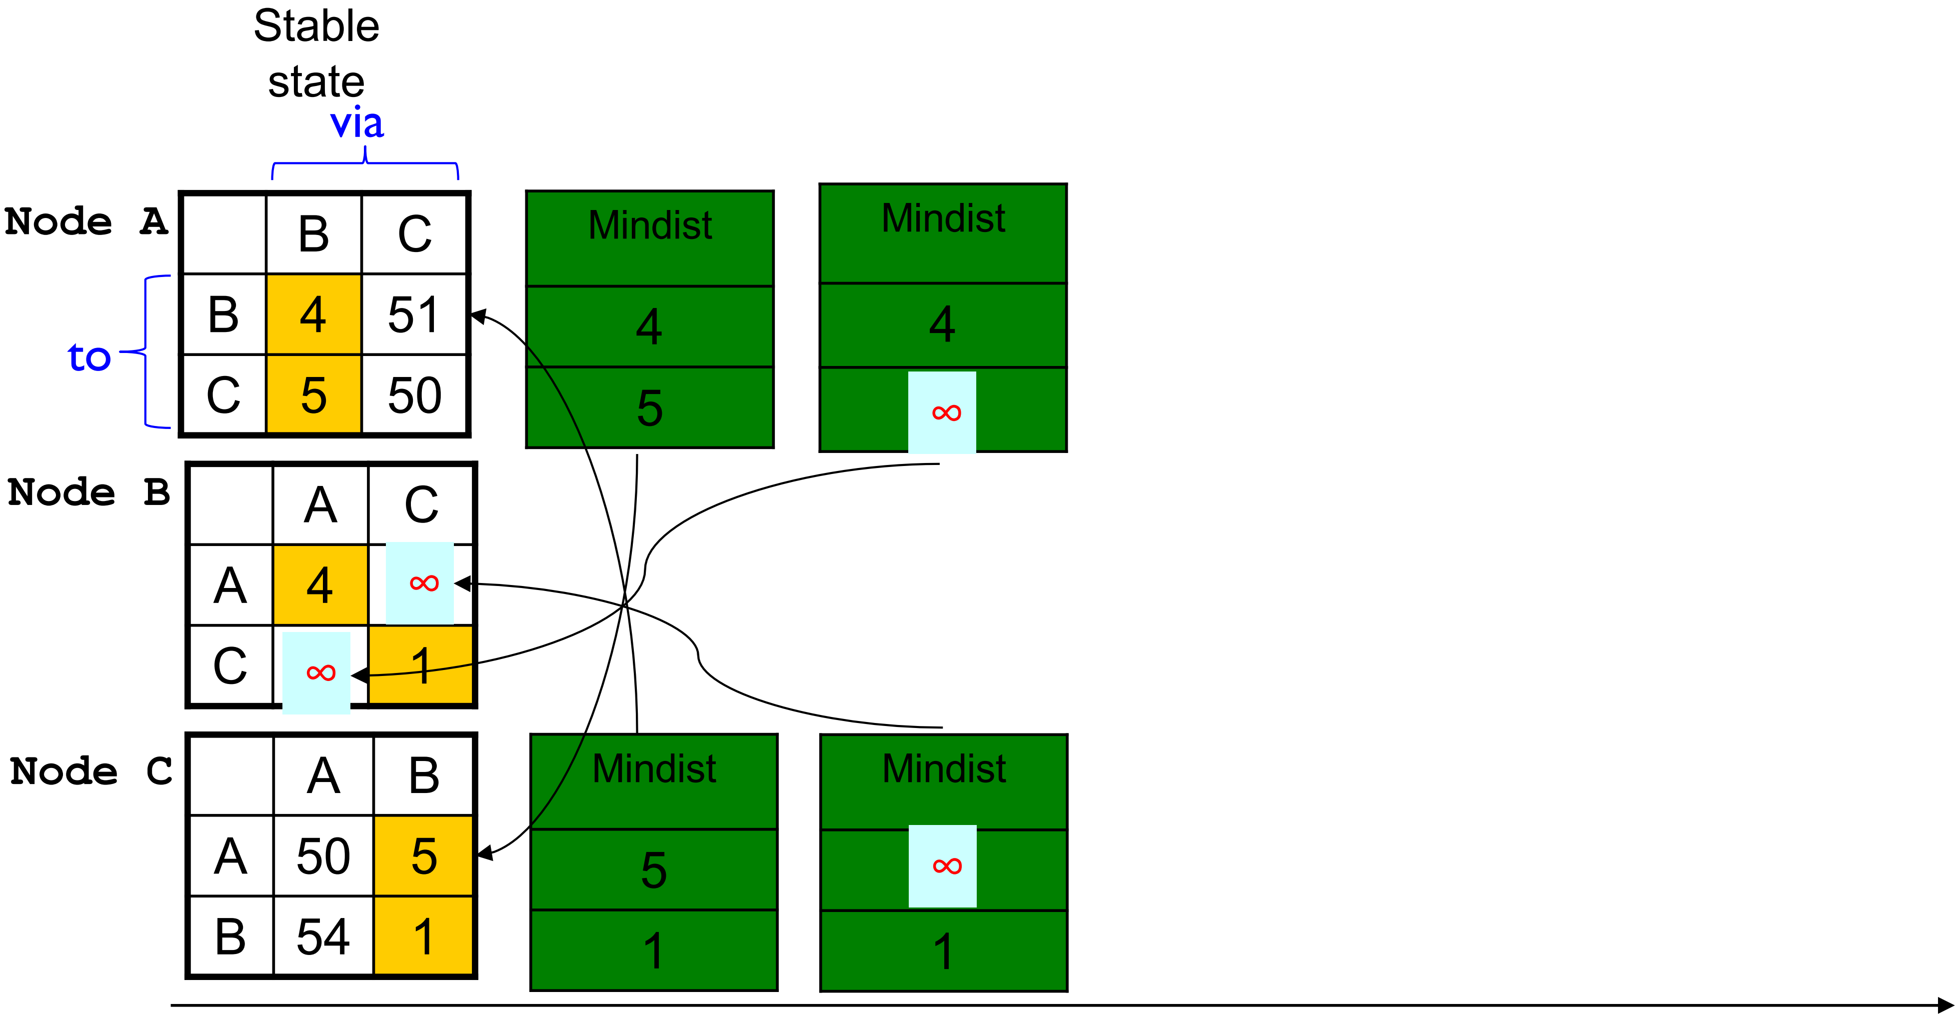

Node A

	B	C
B	4	51
C	5	50

Node B

	A	C
A	4	∞
C	∞	1

Node C

	A	B
A	50	5
B	54	1

Mindist
4
5

Mindist
4
∞

Mindist
5
1

Mindist
∞
1

DV: Poisoned Reverse

If B routes through C to get to A:
B tells C its (B's) distance to A is infinite

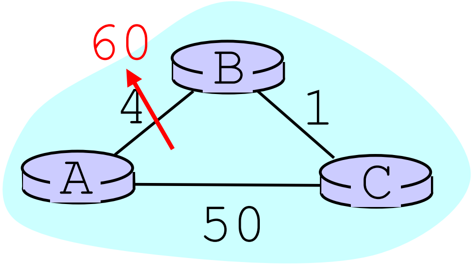

Stable state

Node A (via: columns; to: rows)

	B	C
B	4	51
C	5	50

Node B

	A	C
A	4	∞
C	∞	1

Node C

	A	B
A	50	5
B	54	1

A-B changed

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	6
C	65	1

Node C

	A	B
A	50	5
B	54	1

Link cost changes here

DV: Poisoned Reverse

If B routes through C to get to A:
B tells C its (B's) distance to A is infinite

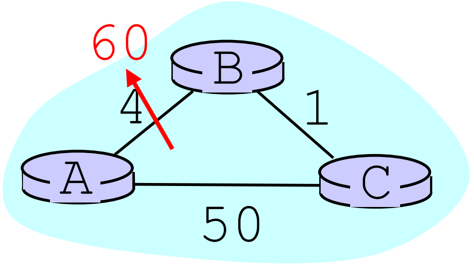

Stable state | A-B changed | A sends its DV to B, C | B sends its DV to A, C

via

to

Node A

Stable state

	B	C
B	4	51
C	5	50

A-B changed

	B	C
B	60	51
C	61	50

A sends its DV to B, C

	B	C
B	60	51
C	61	50

B sends its DV to A, C

	B	C
B	60	51
C	61	50

Node B

Stable state

	A	C
A	4	∞
C	∞	1

A-B changed

	A	C
A	60	∞
C	∞	1

A sends its DV to B, C

	A	C
A	60	∞
C	110	1

B sends its DV to A, C

	A	C
A	60	∞
C	110	1

Node C

Stable state

	A	B
A	50	5
B	54	1

A-B changed

	A	B
A	50	5
B	54	1

A sends its DV to B, C

	A	B
A	50	5
B	∞	1

B sends its DV to A, C

	A	B
A	50	7
B	∞	1

Link cost changes here

DV: Poisoned Reverse

If B routes through C to get to A:
B tells C its (B's) distance to A is infinite

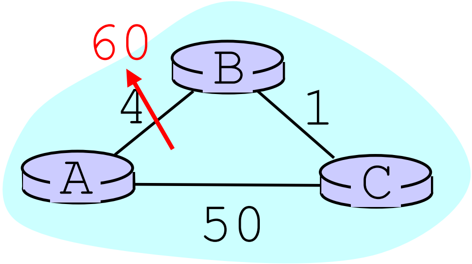

Stable state

via

to

Node A

	B	C
B	4	51
C	5	50

Node B

	A	C
A	4	∞
C	∞	1

Node C

	A	B
A	50	5
B	54	1

A-B changed

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	∞
C	∞	1

Node C

	A	B
A	50	5
B	54	1

A sends its DV to B, C

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	∞
C	110	1

Node C

	A	B
A	50	5
B	∞ |	1

B sends its DV to A, C

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	∞
C	110	1

Node C

	A	B
A	50	61
B	∞ |	1

Link cost changes here

DV: Poisoned Reverse

If B routes through C to get to A:
B tells C its (B's) distance to A is infinite

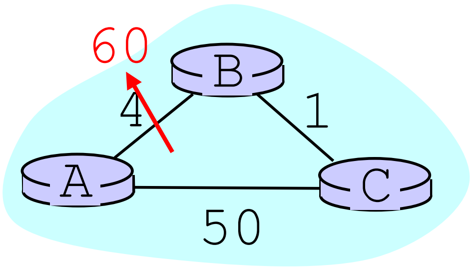

Stable state

via

to

Node A

	B	C
B	4	51
C	5	50

Node B

	A	C
A	4	∞
C	∞	1

Node C

	A	B
A	50	5
B	54	1

A-B changed

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	∞
C	∞	1

Node C

	A	B
A	50	5
B	54	1

A sends its DV to B, C

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	∞
C	110	1

Node C

	A	B
A	50	5
B	∞	1

B sends its DV to A, C

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	∞
C	110	1

Node C

	A	B
A	50	61
B	∞	1

C sends its DV to A, B

Node A

	B	C
B	60	51
C	61	50

Node B

	A	C
A	60	51
C	110	1

Node C

	A	B
A	50	∞
B	∞	1

Link cost changes here

Converges after C receives another update from B

Will Poison-Reverse Completely Solve the Count-to-Infinity Problem?

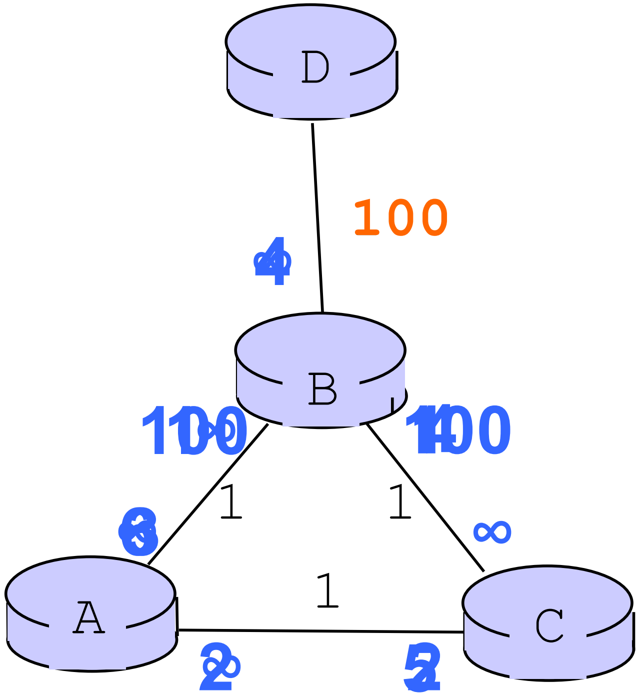

Numbers in blue denote the best cost
to destination D advertised along the link

Quiz: Link-state routing

- In link state routing, each node sends information of its direct links (i.e., link state) to __________?

A. Immediate neighbours
B. All nodes in the network
C. Any one neighbor
D. No one

Quiz: Distance-vector routing

- In distance vector routing, each node shares its distance table with __________?

A. Immediate neighbours
B. All nodes in the network
C. Any one neighbor
D. No one

Quiz: Distance-vector routing

- Which of the following is true of distance vector routing?

A. Convergence delay depends on the topology (nodes and links) and link weights
B. Convergence delay depends on the number of nodes and links
C. Each node knows the entire topology
D. A and C
E. B and C

Each link does't know the entire topology. For instead, they just know the ssp to each node

Comparison of LS and DV algorithms

message complexity

- ***LS:*** with n nodes, E links, O(nE) msgs sent
- ***DV:*** exchange between neighbors only
  - convergence time varies

speed of convergence

- ***LS:*** O(n^2) algorithm requires O(nE) msgs
  - may have oscillations
- ***DV:*** convergence time varies
  - may be routing loops
  - count-to-infinity problem

robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect *link* cost
- each node computes only its *own* table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
  - error propagate thru network

Real Protocols

Link State

Open Shortest Path First (OSPF)

Intermediate system to intermediate system (IS-IS)

Distance Vector

Routing Information Protocol (RIP)

Interior Gateway Routing Protocol (IGRP-Cisco)

Border Gateway Protocol (BGP)

Network layer, control plane: outline

5.1 introduction

5.2 routing protocols

- link state
- distance vector
- hierarchical routing

Self study

5.6 ICMP: The Internet Control Message Protocol

ICMP: Internet Control Message Protocol

- Used by hosts & routers to communicate network level infromation
  - Error reporting: unreachable host, network, port
  - Echo request/reply (used by ping)
- Works above IP layer
  - ICMP messages carried in IP datagrams
- ICMP message: type, code plus IP header and first 8 bytes of IP datagram payload causing error

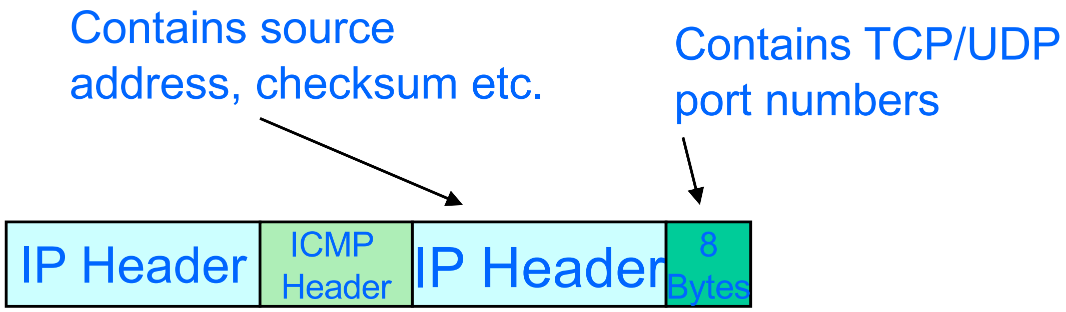

ICMP: Internet Control Message Protocol

Type	Code	Description
0	0	echo reply(ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	3	dest port unreachable
3	4	frag needed; DF set
8	0	echo request(ping)
11	0	TTL expired
11	1	frag reassembly time exceeded
12	0	bad IP header

Traceroute and ICMP

- Source sends series of UDP segments to dest
  - first set has TTL =1
  - second set has TTL=2, etc.
  - unlikely port number
- When *n*th set of datagrams arrives to nth router:
  - router discards datagrams
  - and sends source ICMP messages (type 11, code 0)
  - ICMP messages includes IP address of router
- when ICMP messages arrives, source records RTTs

stopping criteria:

- UDP segment eventually arrives at destination host
- destination returns ICMP "port unreachable" message (type 3, code 3)
- source stops